

TCP-Fuzz: Detecting Memory and Semantic Bugs in TCP Stacks with Fuzzing

Yong-Hao Zou and Jia-Ju Bai, Tsinghua University; Jielong
Zhou, Jianfeng Tan,
and Chenggang Qin, Ant Group; Shi-Min Hu, Tsinghua
University


Background

➤ TCP Stack

- The transport layer protocol of the Internet(or **implementation**).

➤ Black box description

- Input
 - Syscall → `bind()`, `socket()`, `listen()`, ...
 - Packets → TCP segment from IP layer
- Output
 - Syscall result → File descriptor from `socket()`, ...
 - Packets → TCP segment to IP layer



Background

➤ Features of TCP Stack

- Dependency between the two inputs

The input sequence of the TCP protocol stack has certain characteristics

- **Syscall - syscall**

`socket()`, `bind()`, `listen()` are often together and in a fixed order

- **Packets - packets**

The segments of the same TCP connection have the same IP address and port

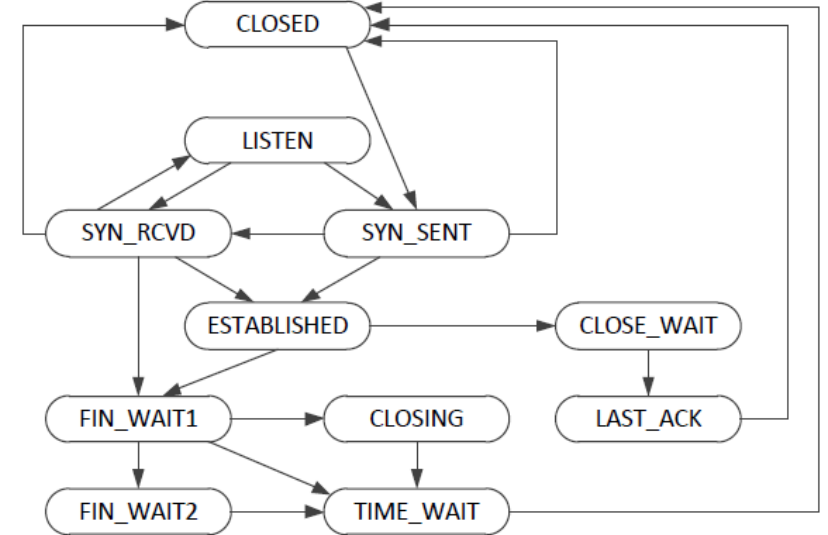
- **Syscall - packets**

`accept()` returns only after the TCP stack receives the last one of the three-way handshake packets.

Background

➤ Features of TCP Stack


- State model
TCP Stack implementations have complex state models.
- Same Semantic rules
The RFC document defines various behaviors of TCP.



Background

➤ Fuzz testing (Fuzzing)

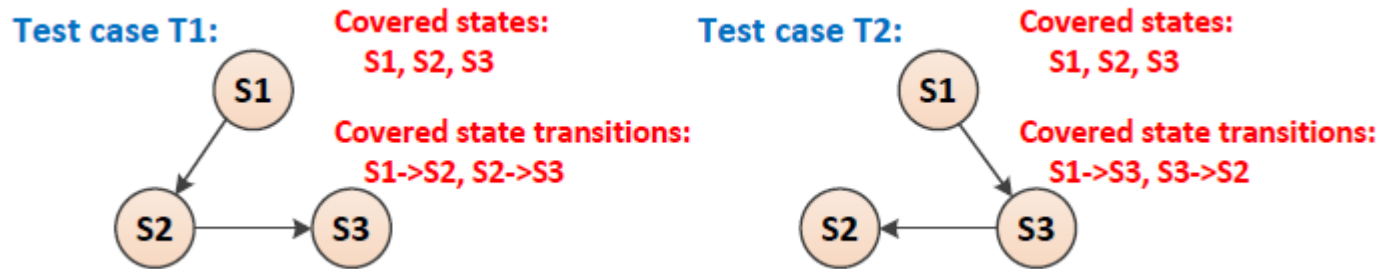
- Automated software testing technique
- Input(test case) generator
 - mutation-based(**need seeds**(input given by user))
 - generation-based(without seeds)
- Application
 - File format
 - **Network protocol**



Problems

➤ Limitations of existing fuzzing for TCP Stack

- Fail to generate two-dimensional inputs with dependencies;
- Neglect the coverage of state transitions.



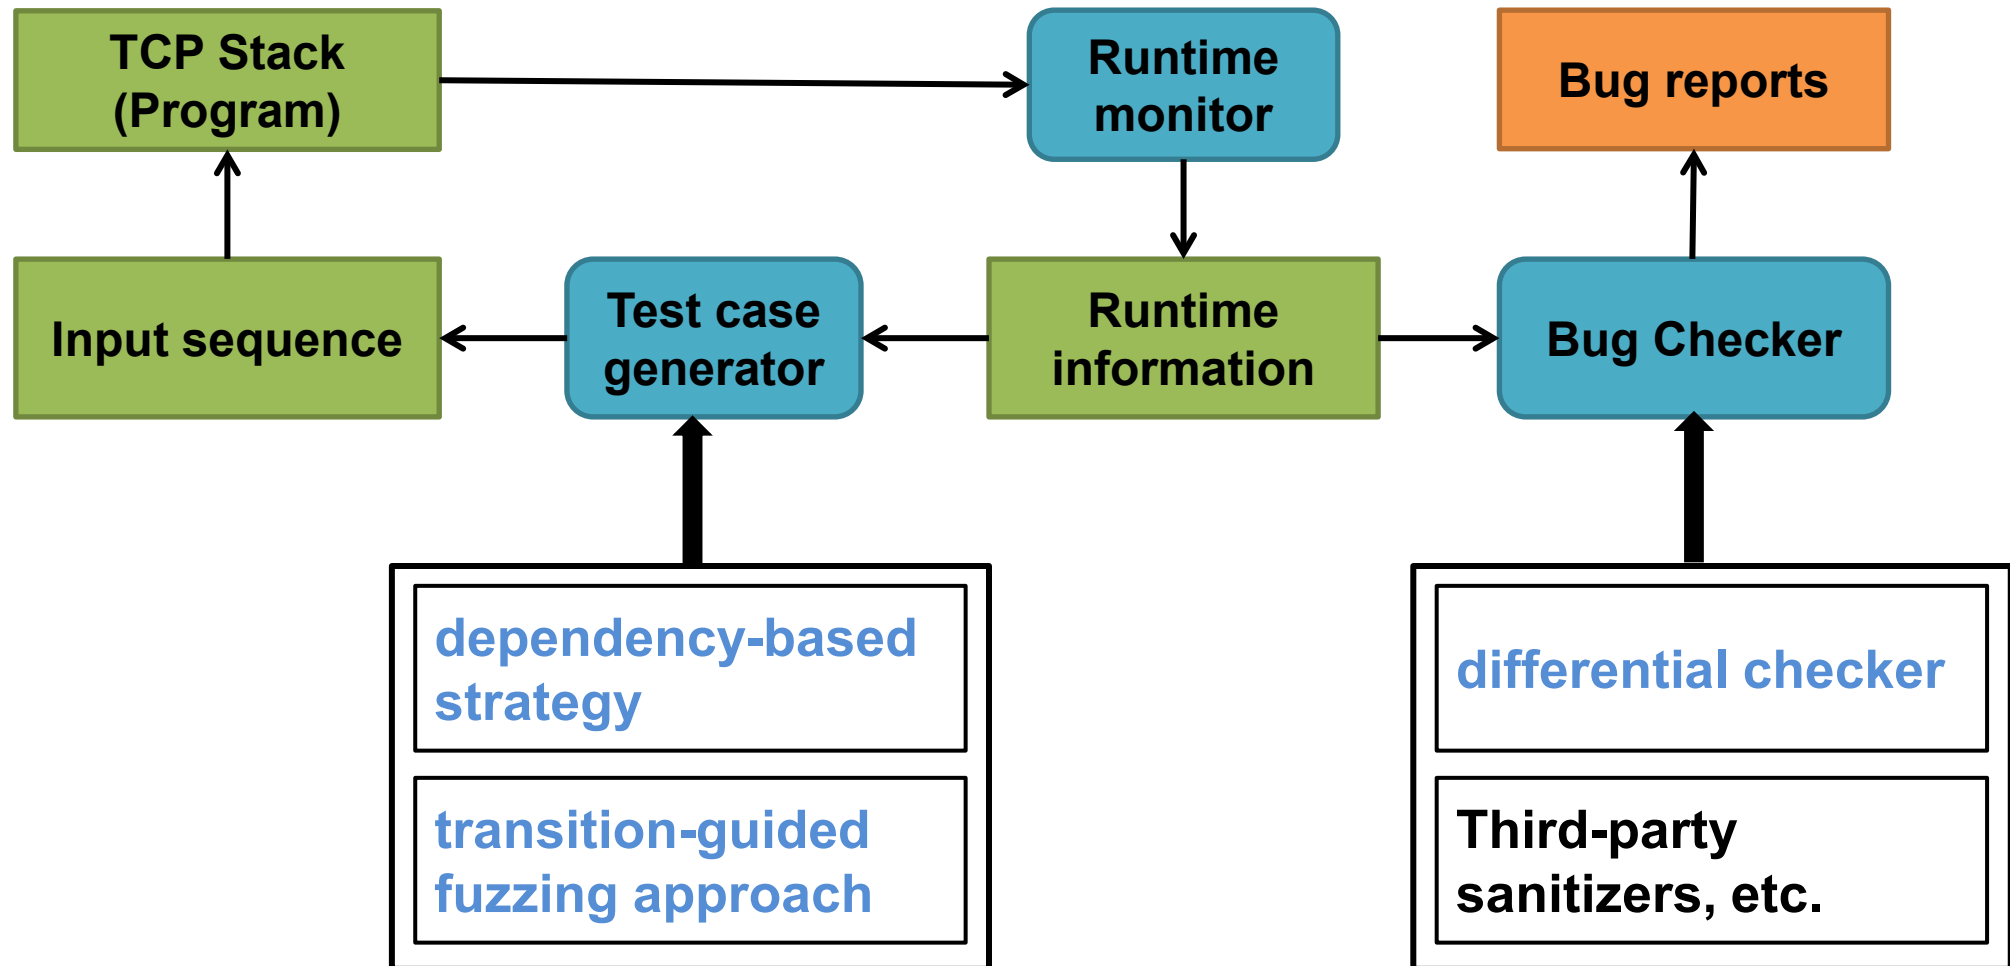
- Lack effective detection of semantic bugs

Time	FreeBSD		mTCP		TLDK	
	<i>Memory</i>	<i>Semantic</i>	<i>Memory</i>	<i>Semantic</i>	<i>Memory</i>	<i>Semantic</i>
2017	2	26	2	6	1	11
2018	9	51	0	4	0	4
2019	9	65	1	3	2	5
Total	20	142	3	13	3	20

Ideas

Features	Problems	Solution
Dependency between the two inputs	Lack of dependency on input generated by the test	dependency-based strategy to generate effective test case
State model	Neglect the coverage of state transitions	transition-guided fuzzing approach to improve the coverage of state transitions
Same Semantic rules	Lack effective detection of semantic bugs	a differential checker to detect semantic bugs

Design



Dependency-Based Strategy

- Discard input sequences that do not meet certain dependency rules;
- Generate sequence according to dependency rules;
 - Dependency rules are defined in advance
- Rule Example:

Kind	Dependency rule
Syscall-syscall(x5)	Socket() and connect() are called in order when a connection is active open.
Packet-packet(x5)	After a connection is established, the source port and destination port of each packet are fixed
Syscall-packet(x5)	After close is called, a packet with the FIN flag should be sent

Transition-Guided Fuzzing Approach

➤ State

- Branch coverage (as existing fuzzing approaches do)
- An example state $s_1 = \{1, 0, 1, 0\}$


BR1	BR2	BR3	BR4
1	0	1	0

➤ State transitions

- State difference $BT = S_{i+1} - S_i$ (new)

Transition-Guided Fuzzing Approach

➤ Example

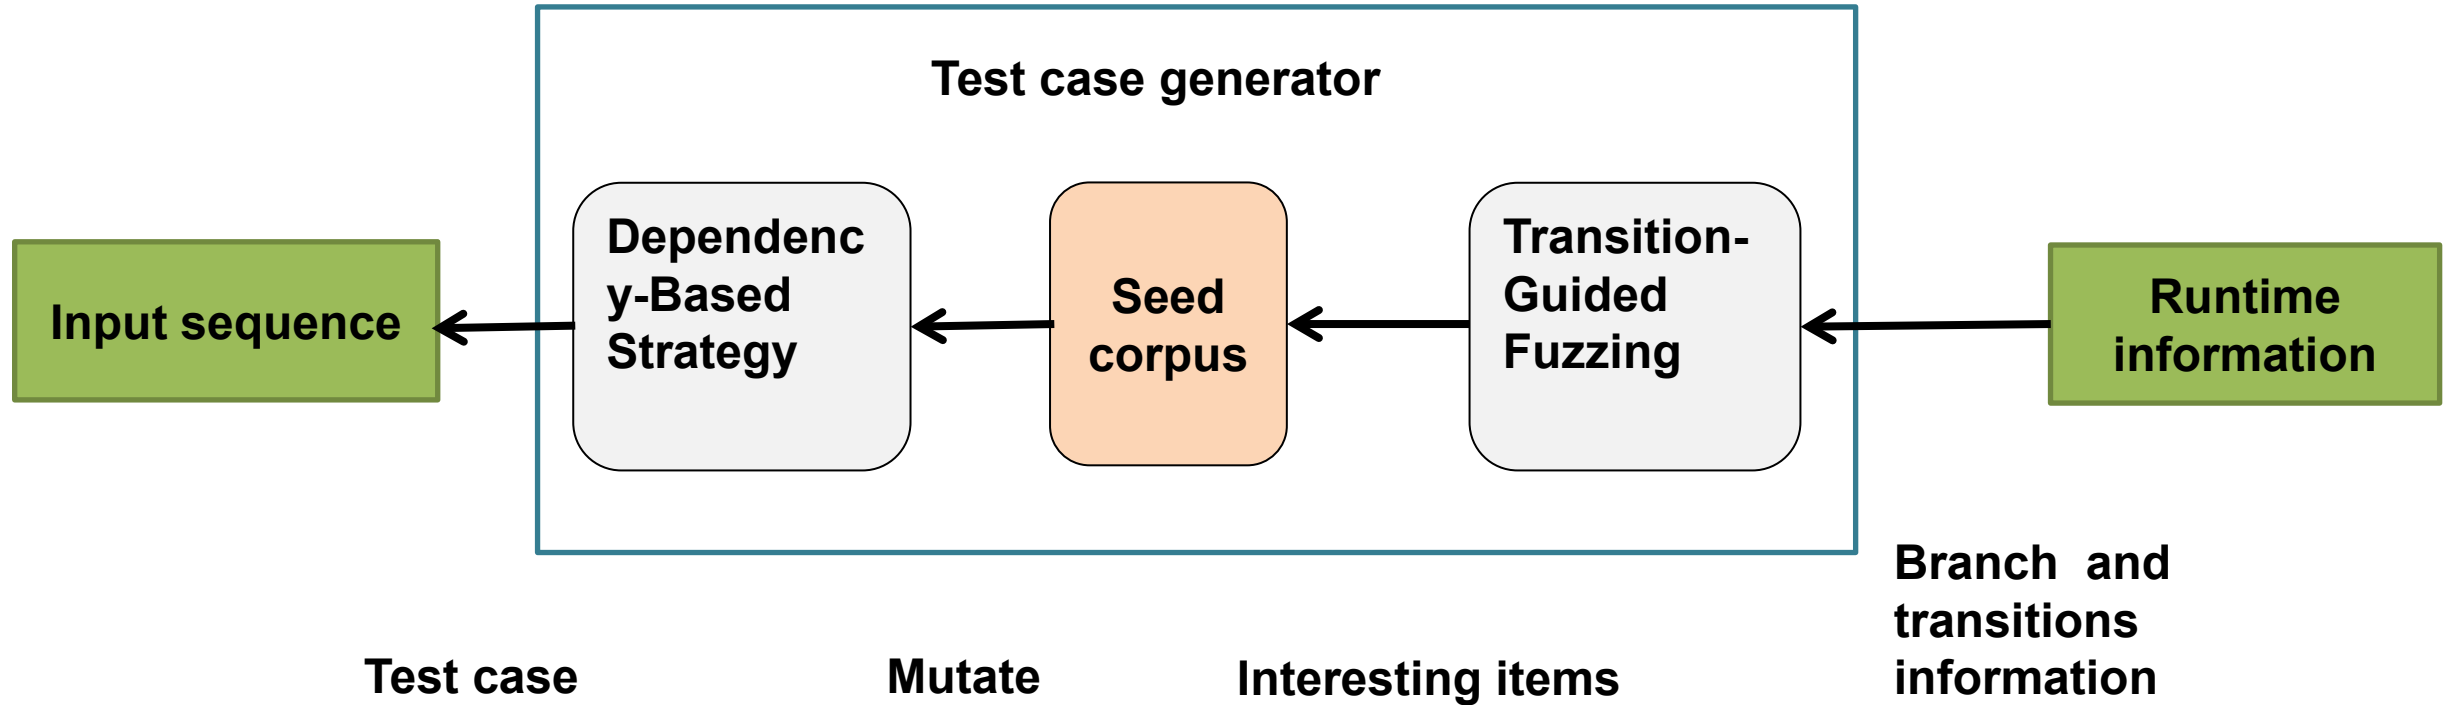


When input item4 :

State coverage only: S1 has appeared → input item4 is **useless**


Transition-guided: T3 is new → Input item4 is **interesting**

Transition-Guided Fuzzing Approach



Differential Checker

- Check the output of the same input on different TCP Stacks
 - Outputs are different → at least one has a bug;
 - Output of the tested TCP stack is different from the standard output → the tested TCP stack has a bug;



Evaluation

➤ Test objects


- user-level
 - TLDK (recent);
 - F-Stack (recent);
 - mTCP (well-known in academic).
- Kernel-level (**test cases generated from the user-level TCP stacks**, also used as the reference TCP Stack)
 - FreeBSD TCP (classical);
 - Linux TCP (classical).

Evaluation

➤ Results (without comparison)

- Covers many more branch transitions than branches;
- Covers few new branches during the later tests.

Stack	Testing coverage		Found bugs	
	<i>Branch</i>	<i>Transition</i>	<i>Memory / Semantic</i>	<i>Confirmed / Fixed</i>
TLDK	1.3K	329.4K	2 / 26	28 / 19
F-Stack	7.5K	46.8K	1 / 6	6 / 1
mTCP	1.2K	47.9K	5 / 9	0 / 0
FreeBSD	-	-	0 / 6	5 / 2
Linux	-	-	0 / 1	1 / 1
Total	10.0K	424.1K	8 / 48	40 / 23




Evaluation


➤ Compare Objects

- AFL-like (code coverage);
- Syzkaller-like (code coverage);
- Boofuzz, Fuzzotron, AFLNet (open-source protocol fuzzing approaches).

➤ Results(with comparison, TLDK as example)



In 26 semantic bugs:
25 of these are found by
differential checker



Conclusion

- Find problems and solve them one by one
 - **Lack of dependency on input generated by the test;**
Dependency-based strategy to generate effective test case
 - **Neglect the coverage of state transitions;**
Transition-guided fuzzing approach to improve the coverage of state transitions
 - **Lack effective detection of semantic bugs.**
Design differential checker to detect semantic bugs