



Zero-Change Object Transmission for Distributed Big Data Analytics


Mingyu Wu, Institute of Parallel and Distributed Systems, SEIEE, Shanghai Jiao Tong University and Shanghai AI Laboratory; Shuaiwei Wang, Institute of Parallel and Distributed Systems, SEIEE, Shanghai Jiao Tong University; Haibo Chen and Binyu Zang, Institute of Parallel and Distributed Systems, SEIEE, Shanghai Jiao Tong University and Engineering Research Center for Domain-specific Operating Systems, Ministry of Education, China



Background


➤ Distributed big data

- Data is distributed in different server
- Data needs to be exchanged between servers during data analysis
- Data is generally read-only



➤ Server runtime environment

- JVM (java virtual machine)
- Data is exchanged between different JVM instances

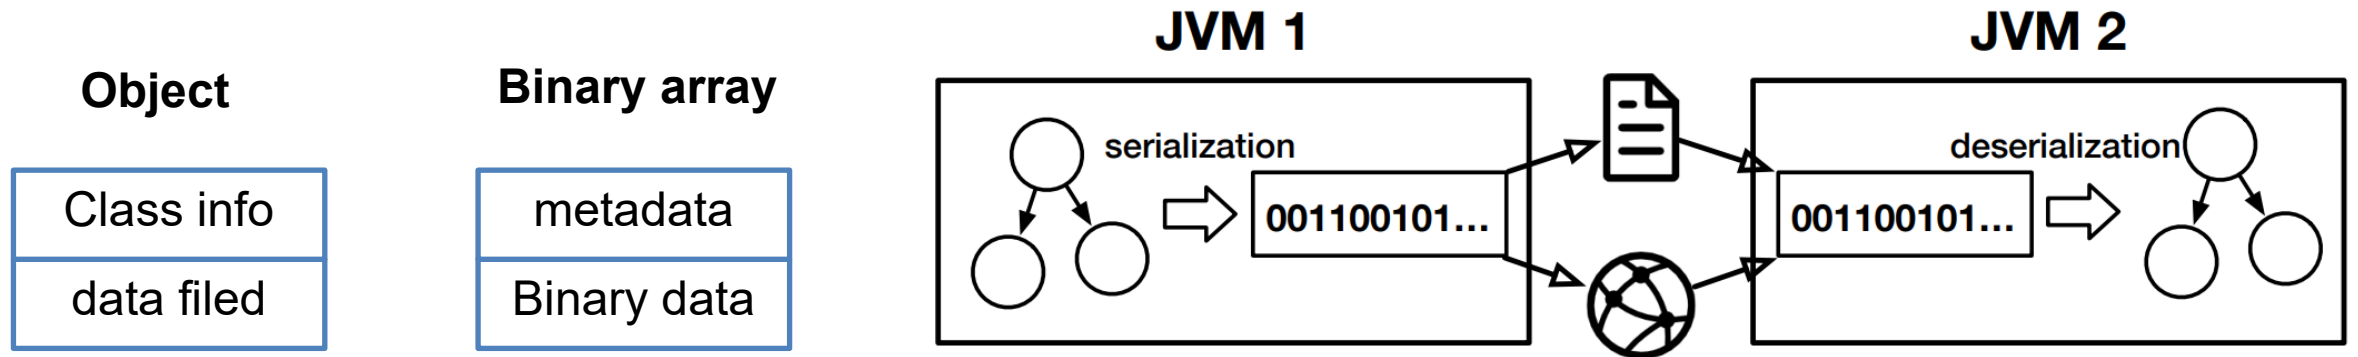




Background

➤ How to exchange: **Serialization and Deserialization**

- Data is stored in JVM's memory in the form of objects
- **Serialization:** Object -> binary array (standardized method)
- **Deserialization:** Binary array -> Object (standardized method)



Problem

➤ CPU overhead

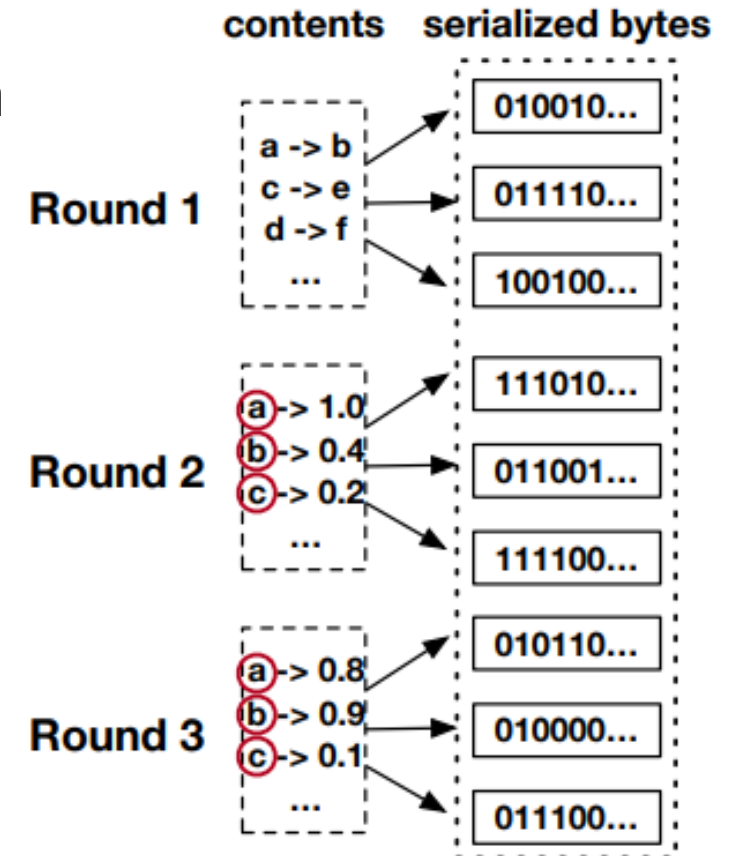
- Object \leftarrow [CPU] \rightarrow Binary array

➤ Memory overhead

- Additional memory overhead during (de)serialization
- Data redundancy due to standardized methods (extra metadata)

➤ Overhead from repeated data transmission (in specific scenario)

- iterative algorithm
- ...



Idea & Challenge


➤ IDEA

Remove (De)serialization

Use distributed shared memory for data exchange

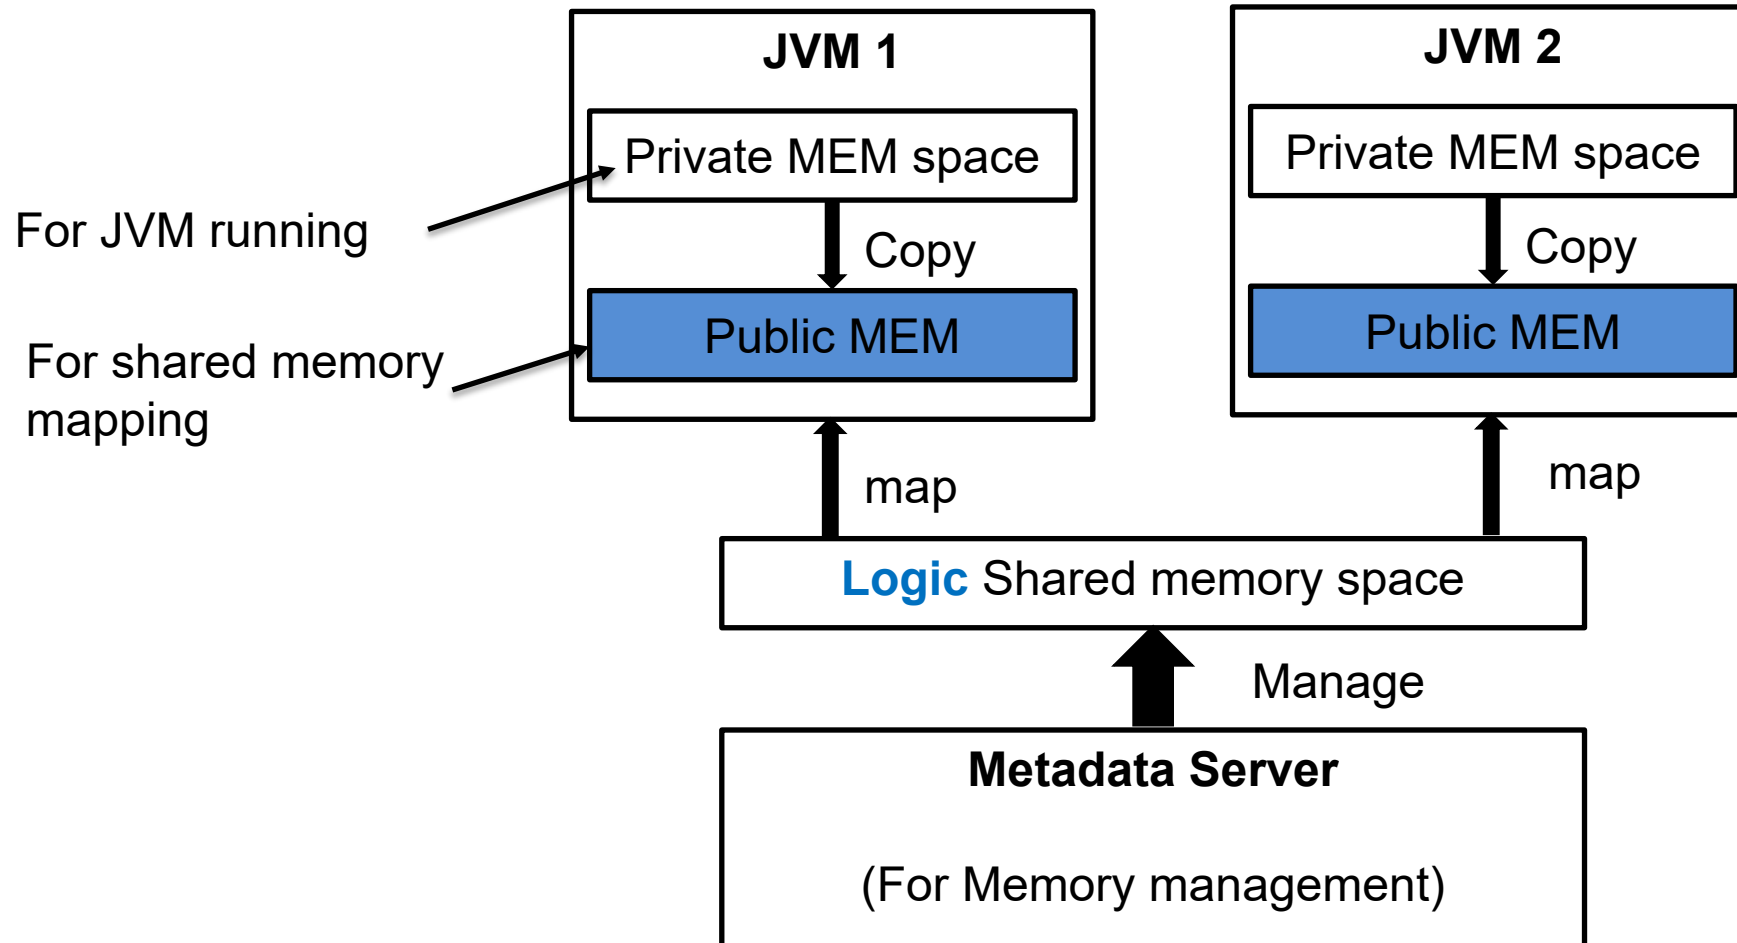
➤ Challenge

- Architecture design
- Data exchange protocol
- Shared memory management
- Remove redundant data transmission



Architecture Design

➤ Architecture





Lazy data exchange protocol

➤ Sender

- Acquire memory from shared memory space
- Copy data to public MEM in specific address (still in local memory)

➤ Receiver

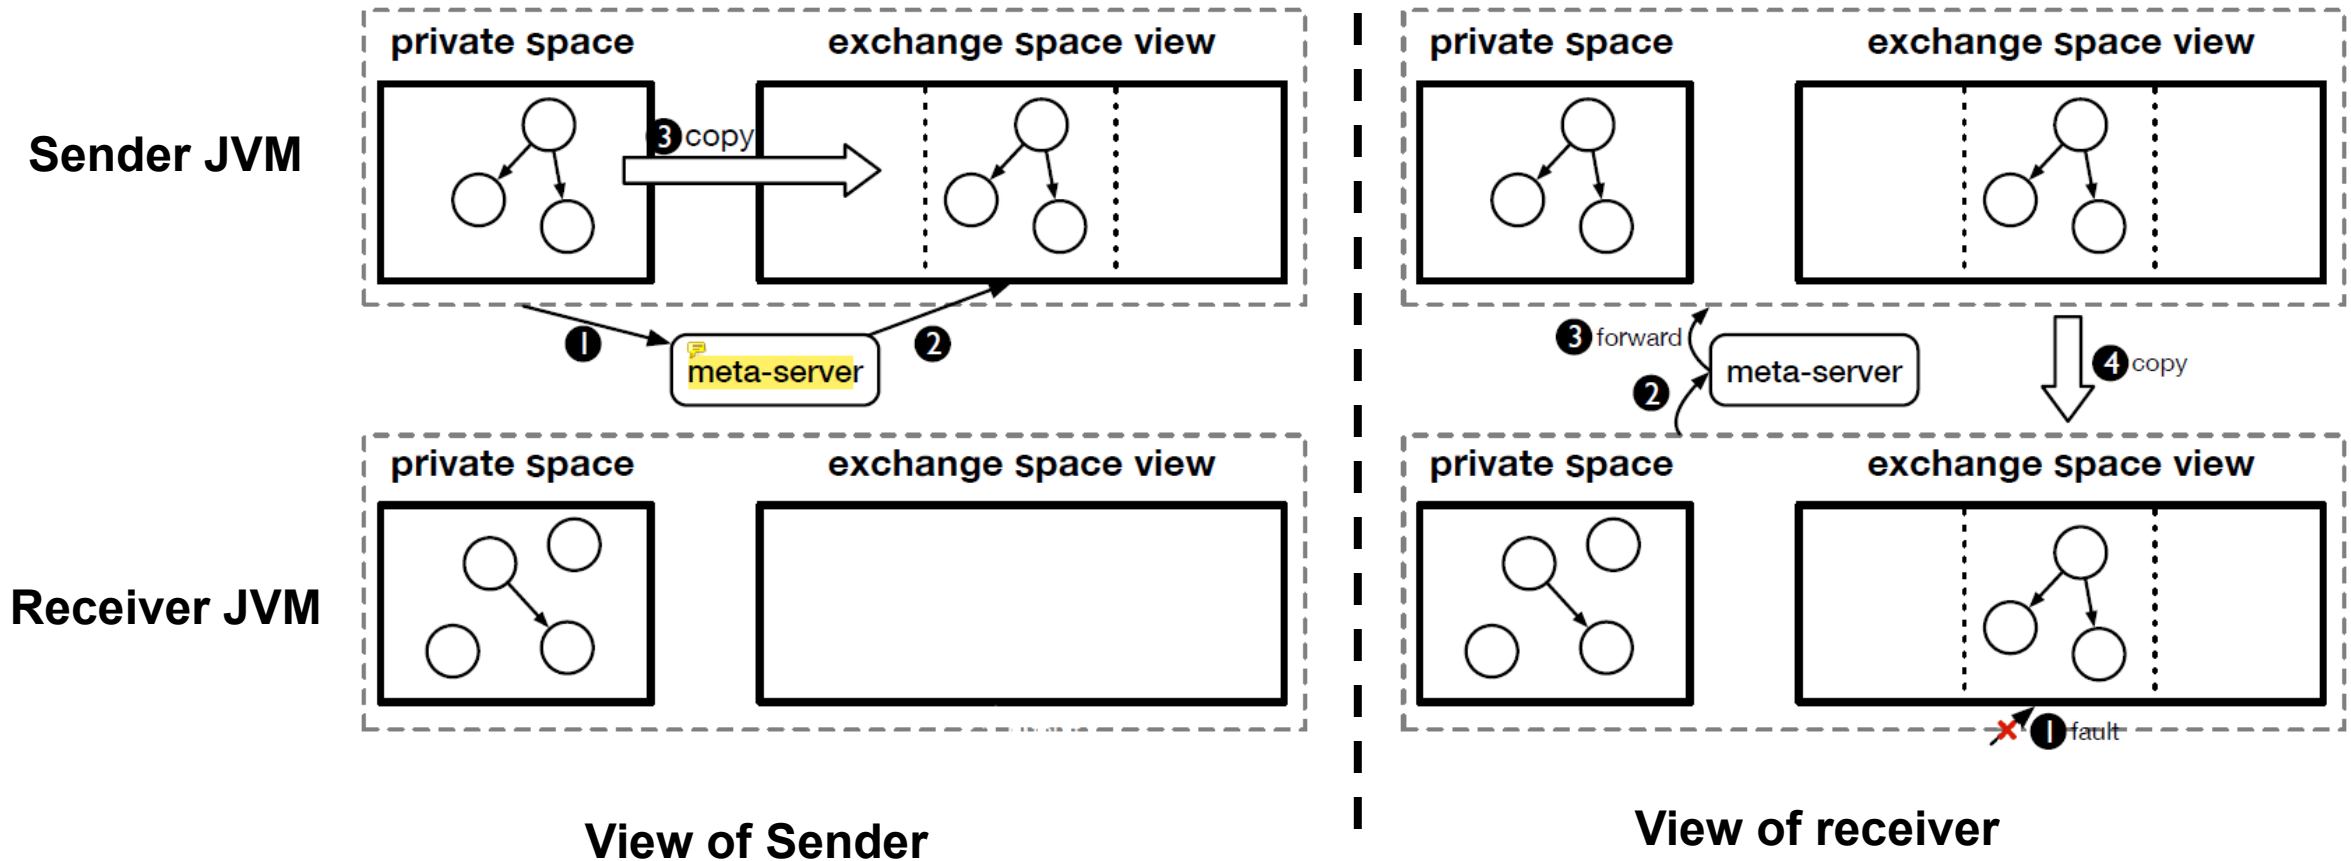
- Read remote memory from shared memory space
- Memory is in local -> Read Directly
- Memory is in remote (mapped to another JVM) -> Build network connection

➤ Problem & Guess : How to know the address of data

- There is additional network communication to exchange extra info

Data exchange protocol

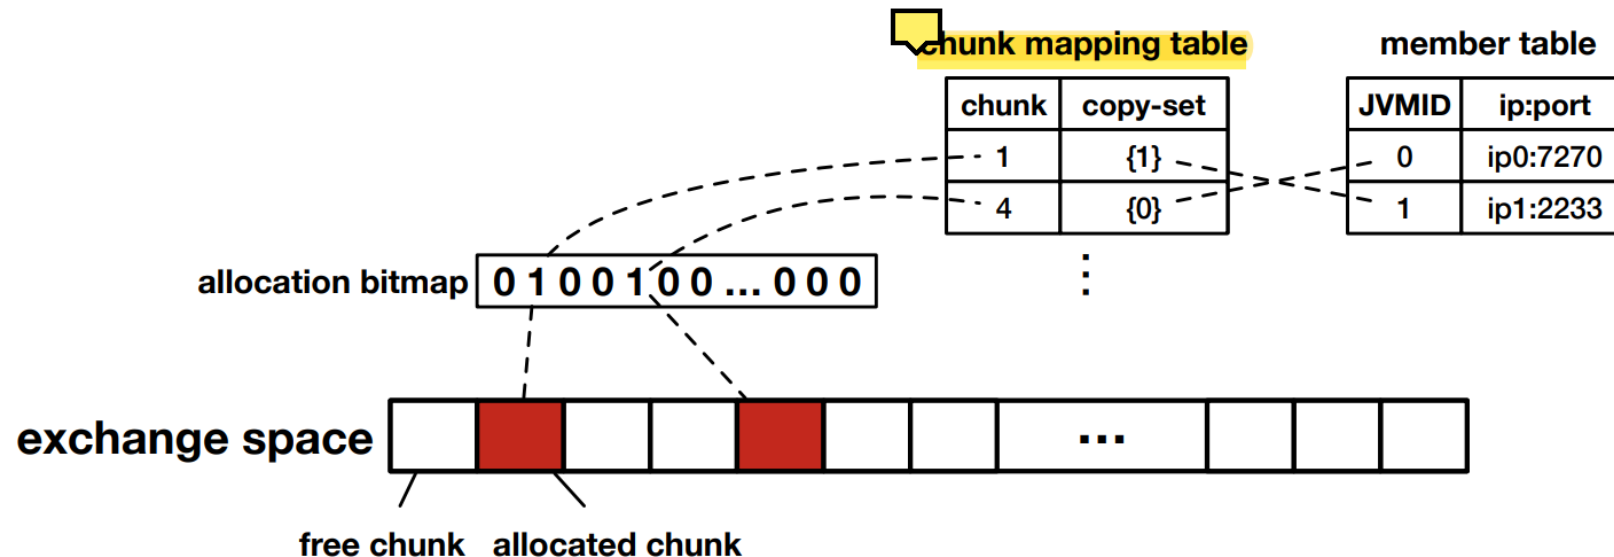
➤ Example



Shared Memory Management (Metadata server)

➤ Logical shared memory layout

- Allocation bitmap
- Member table (JVM info table)
- Chunk mapping table (chunk id → {JVMs})
(check if a JVM has a copy of this memory chunk)





Shared Memory Management (JVM)

➤ **Acquire()->address**

- When: has no public memory to store data
- How: Metadata server: scan bitmap and allocate chunk

➤ **Get remote(address)**

- When: memory needed is not in local
- How: Metadata server find memory location and help build connection to copy memory chunk

➤ **Release(address)**

- When: memory chunk is not needed (or GC)
- How: Metadata server: reclaim memory

Question: How JVM manage their public memory? (no answer)

Data Duplication in Exchange(?)

➤ Build reference between memory chunks

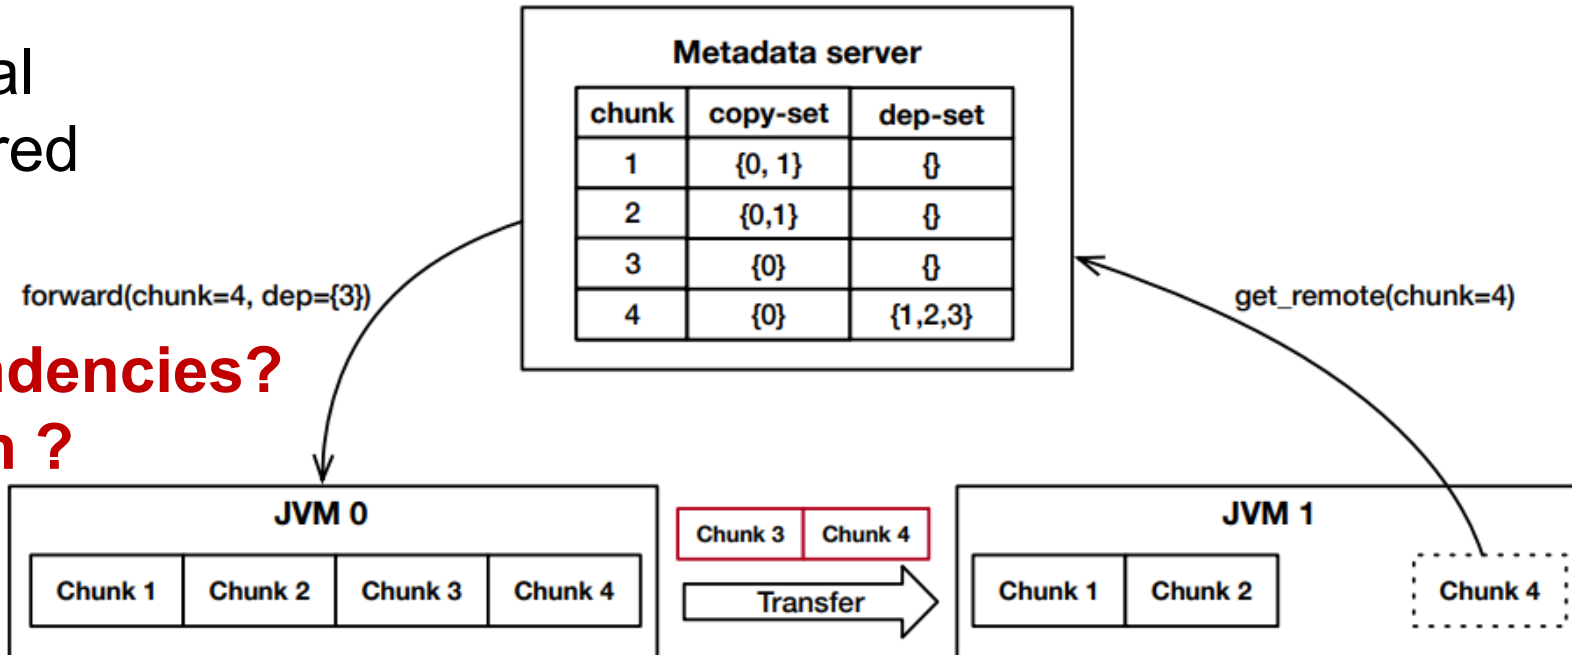
- B reference A : Some objects in A are the same as those in B

➤ Only dependent chunks will be transferred

1. JVM: need chunk 4
2. 4 \rightarrow {1,2,3}
3. 1,2 are already in local
4. Only 3 will be transferred

➤ Problems

- How to maintain dependencies?
- How to save bandwidth ?





Evaluation

➤ Hardware

- Cluster with 4 nodes + Xeon E5-2650 CPU * 2 + 128GB DRAM

➤ Software platform

- JVM: Hotspot JVM 11
- Microbenchmark
- Spark (data analytics engine)
- Flink (data processing engine)
- Compare objects
 - JSR → (standard method, baseline)
 - Kryo → (refined binary data format, smaller size)
 - Naos → (bases on RDMA, no memory copy/ (De)serialization)
 - Skyway → (no (De)serialization + extra metadata)
 - **ZCOT**

Evaluation

➤ Microbenchmark

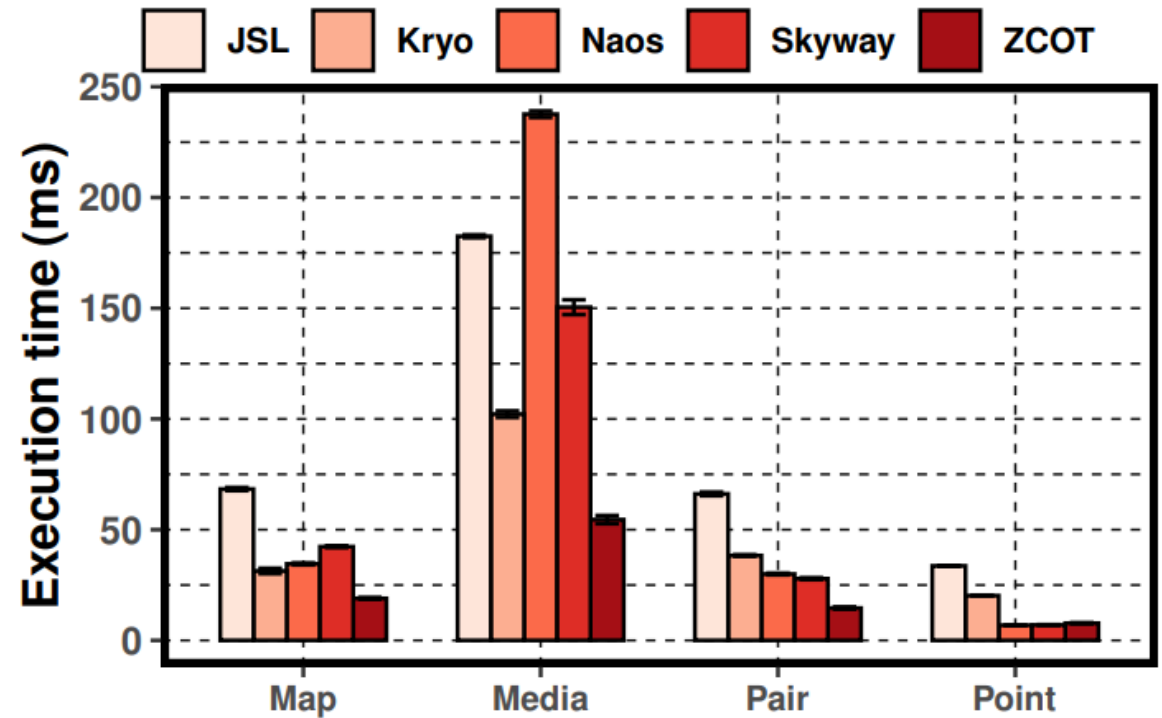
Under two machines

➤ Overall time

Media > Map > Pair > Point
(different data complexities)

➤ Performance

ZCOT > (JSL, Kryo, Naos, Skyway)



Evaluation

➤ Spark

➤ Overall

ZCOT > Kryo > JSL

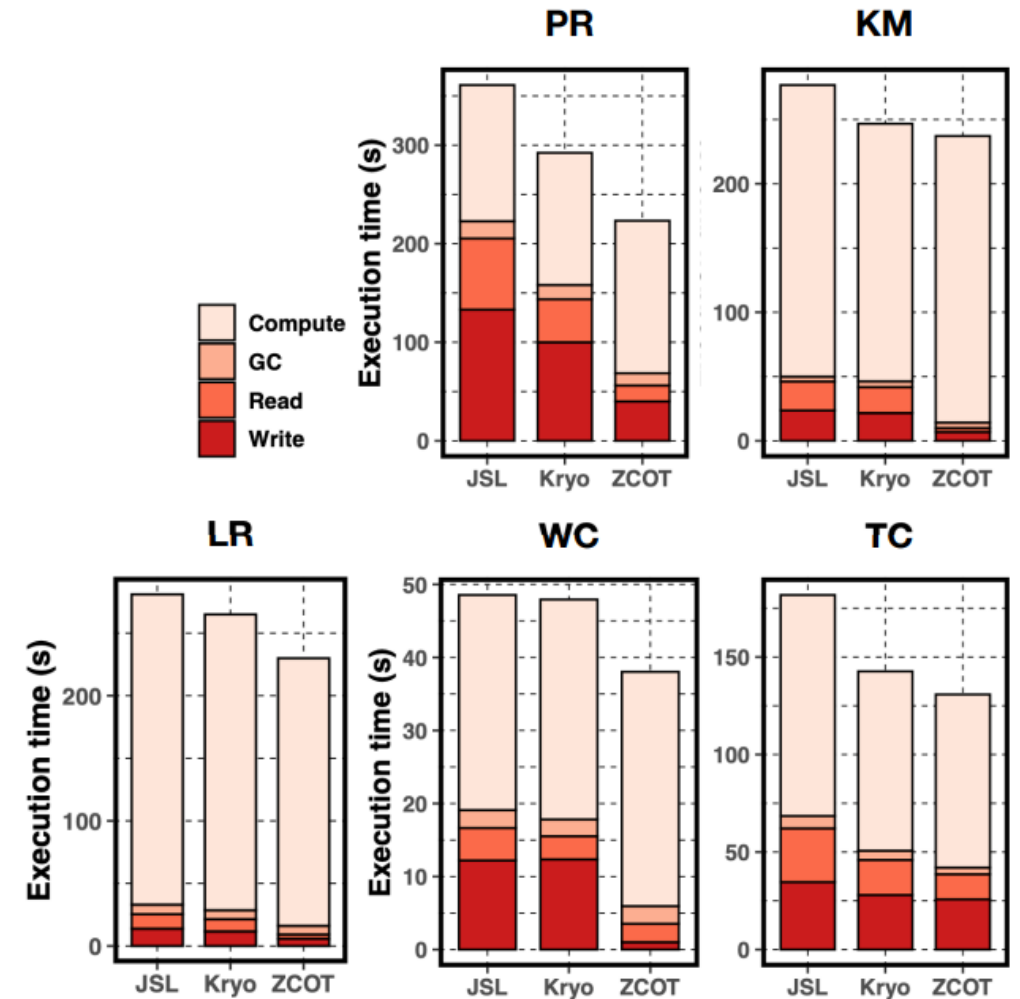
(Lower read and write times)

➤ Total data transferred

Data de-duplication works

	PR	WC	TC	KM	LR
dedup	15.25	4.13	5.03	5.37	5.55
no-dedup	31.64	5.50	10.88	5.86	6.04

Application	Dataset
PageRank (PR)	LiveJournal [4]
Word Count (WC)	LiveJournal
KMeans (KM)	USCensus1990 [10]
Transitive Closure (TC)	Blogs [1, 17]
Logistic Regression (LR)	SUSY [5]





Conclusion

Problem

(De)Serialization overhead In Big Data Analytics

Idea

Use **distributed shared memory** to remove serialization process

Design

DSM design & management

Data exchange
protocol

Remove duplicate data
transmission