

DEPART: Replica Decoupling for Distributed Key-Value Storage

Qiang Zhang and Yongkun Li, University of Science and Technology of China; Patrick P. C. Lee, The Chinese University of Hong Kong; Yinlong Xu, Anhui Province Key Laboratory of High Performance Computing, University of Science and Technology of China; Si Wu, University of Science and Technology of China

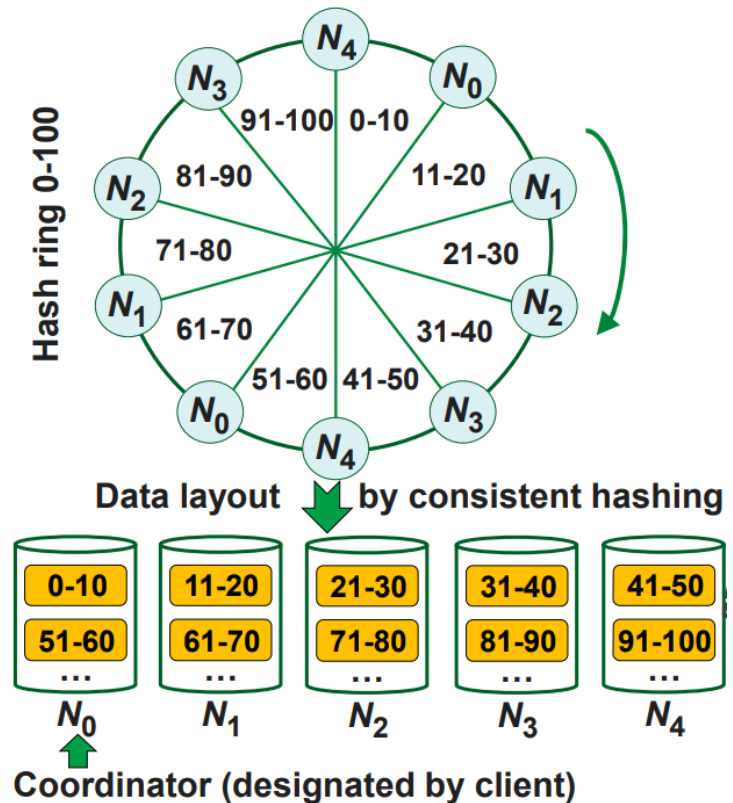

Background

➤ Distributed Storage Server

- Multiple storage nodes
- Consistent hash ring

➤ Client behavior

- Calculate consistent hashes and find nodes




Background

➤ Fault tolerance & Reliability

- Large systems have the potential for error
- Data loss or inconsistency

➤ Solution: Replication

- Each KV-pair is stored in multiple nodes
- Two kinds of data: primary and redundant



Background

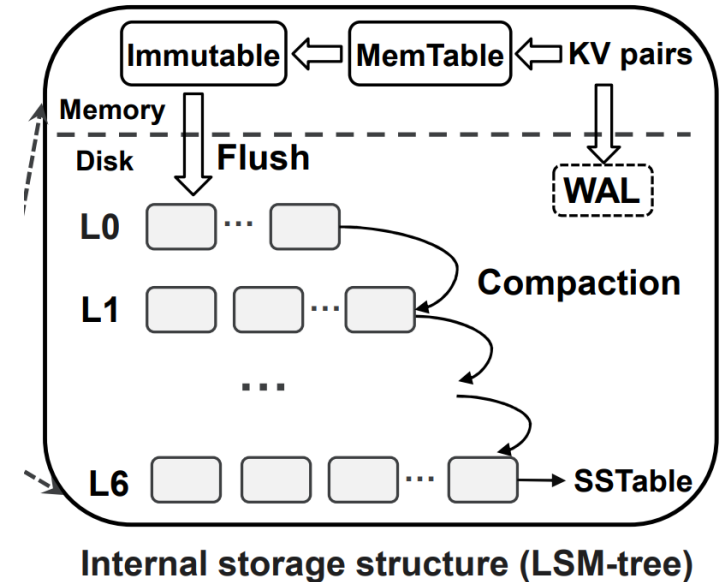
➤ In-node Storage structure: LSM-Tree

- MEM(Tree) + Disk
- Multi-Level tree (KV-pairs is **sorted** by key in each level)
- The capacity of each level is limited;
the higher the level, the bigger the capacity
- Write: Mem table → Immutable
-> Level -> **Compaction**(if full)
- Read: Query top level to bottom level

Write amplification

Read amplification

The larger the LSM-tree, the more serious the read/write amplification



Problem


➤ Uniform indexing

Primary data and redundant data are stored in the same LSM tree(Cassandra, TiKV, etc...)


➤ Problem

Uniform indexing cause serious read amplification and write amplification

Verification




(a) Impact of replication in Cassandra and TiKV

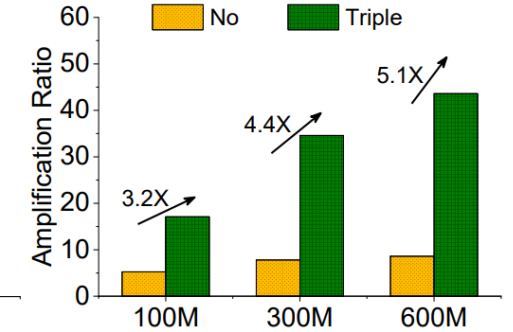


(b) Impact of KV store size in Cassandra

Write 300 M KV-pairs



(a) Impact of replication in Cassandra and TiKV



(b) Impact of KV store size in Cassandra

Read from 300 M

Increase degree of replication

Increase KV store size



**Increase
LSM-tree size**



**read/write
amplification**

Basic Idea

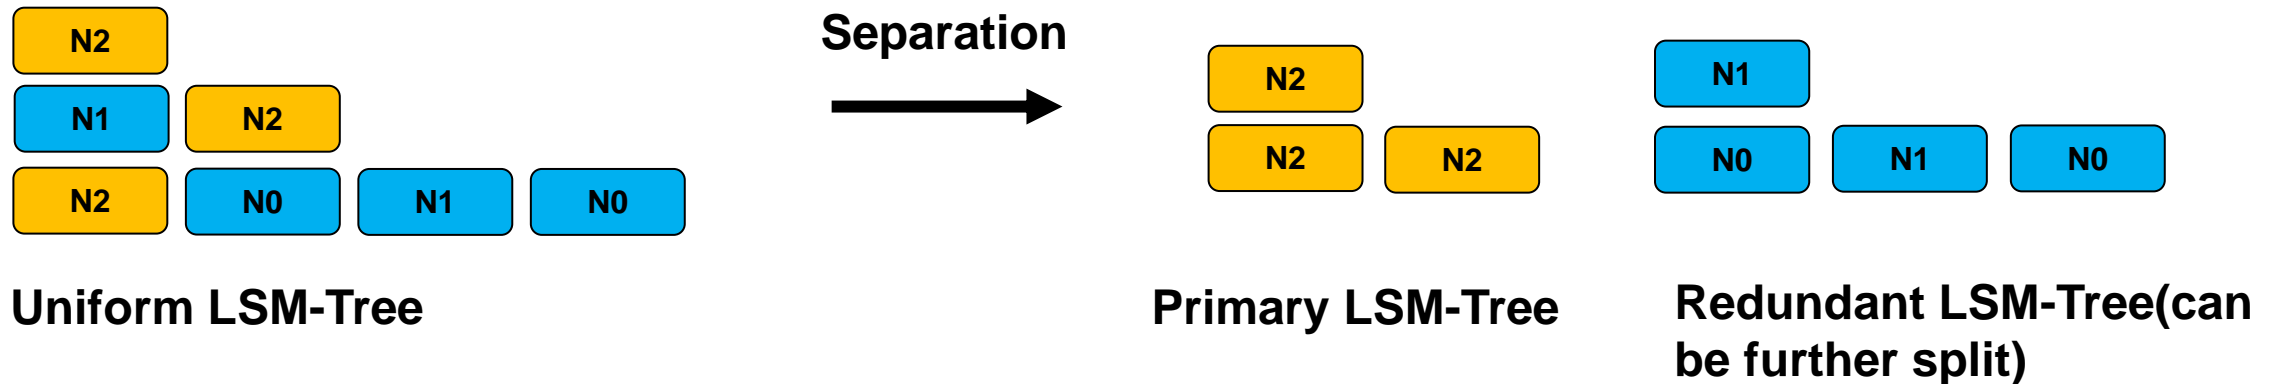
➤ Idea

- **Separate** primary and redundant data to reduce LSM-tree size

Plain Design

➤ Multiple LSM trees

- Separate uniform tree into primary data tree and redundant tree




➤ Evaluation

- Advantage : Reduce read/write amplification
- Problem: More memory overhead(Same memory usage for each tree)


Idea

➤ Consistency

- Read/Write consistency level (RCL/WCL)
 - The minimum number of successful reads/writes that need to be confirmed for a successful operation
- Consistency level is a Balance on Performance and Reliability
- Quantified consistency Level : $WCL + RCL$



Success R/W operation
in $RCL/WCL = 1$



Fixed consistency Level

Idea

➤ LSM-Tree

- Strong Sorted in each level(High read performance, low write performance) **Not suitable for $WCL > RCL$**

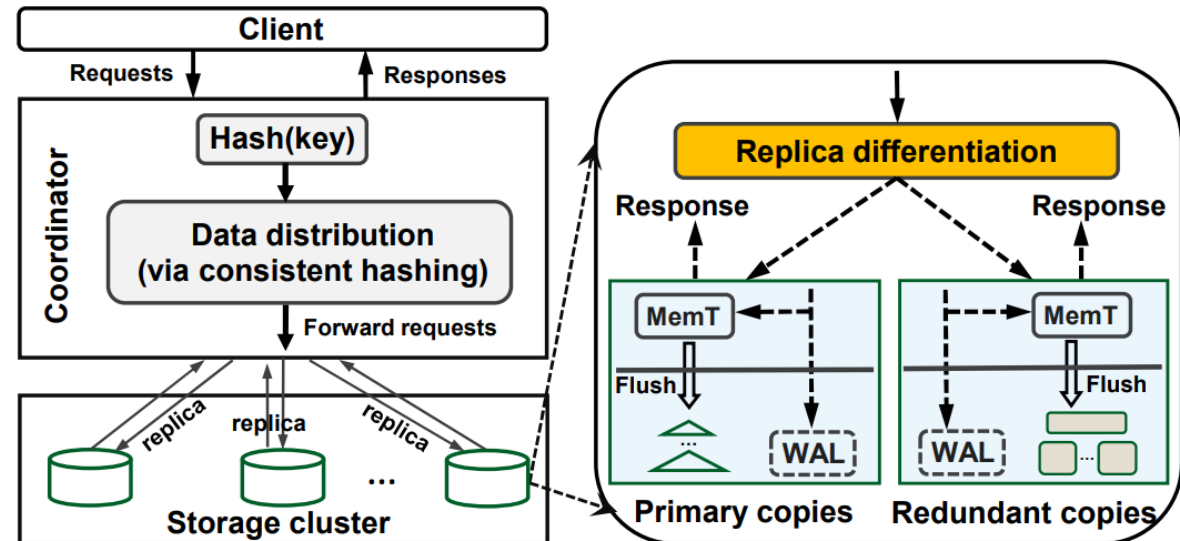
➤ Idea

- **Trick on consistency level : Balance on Read and Write**
 - $WCL > RCL$ -> Guarantee write performance -> weakly ordered
 - $WCL < RCL$ -> Guarantee read performance -> strongly ordered

Design

➤ Separation + Two-layer log

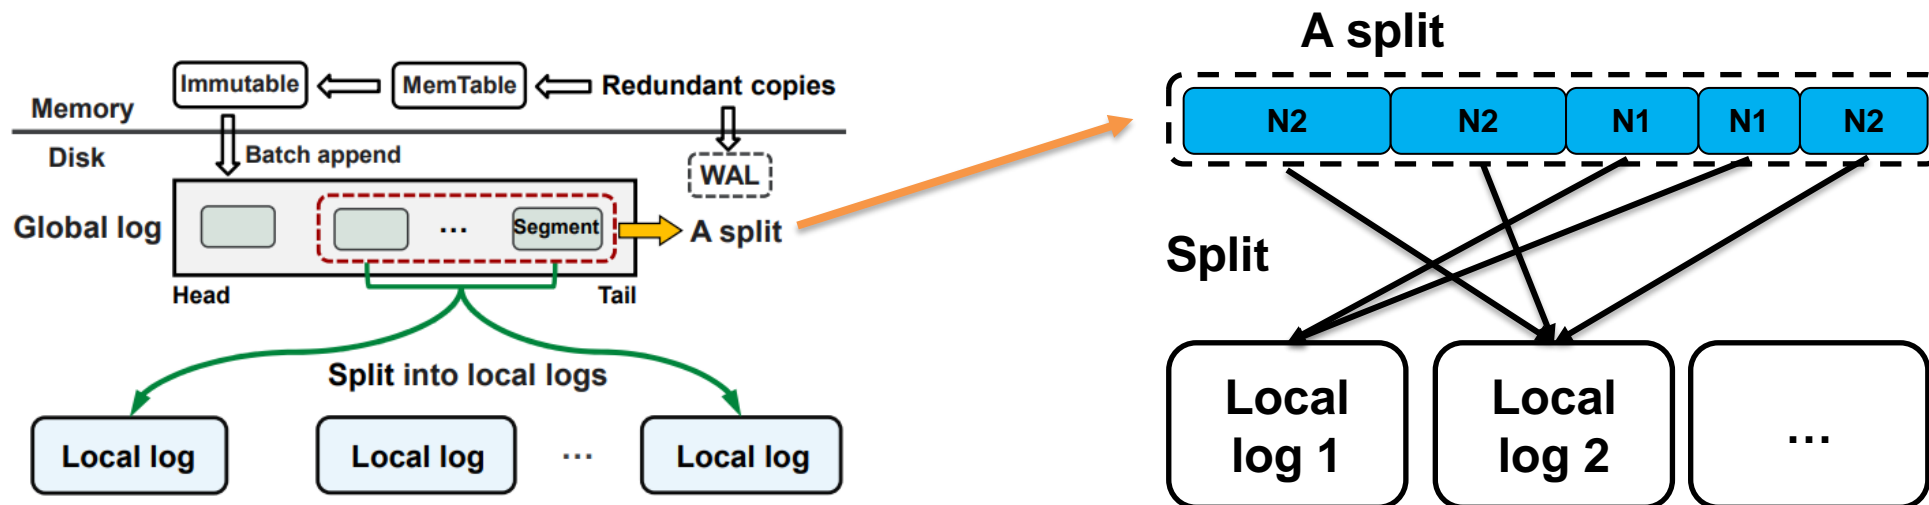
- Replica differentiation : recalculate hash to check if a KV pair is primary or redundant
- Primary data : LSM-Tree
- Redundant data : Two-layer log
 - global log
 - local log



Design

➤ Two-layer log

- Global log : save all flushed KV-pairs without any extra metadata
 - Feature Inline + Append-only → **Fast write**
 - Problem **Hard to recovery(read) + garbage collection cost**
- Local log : take some log from the end of the global log and save them to local log (**background**)







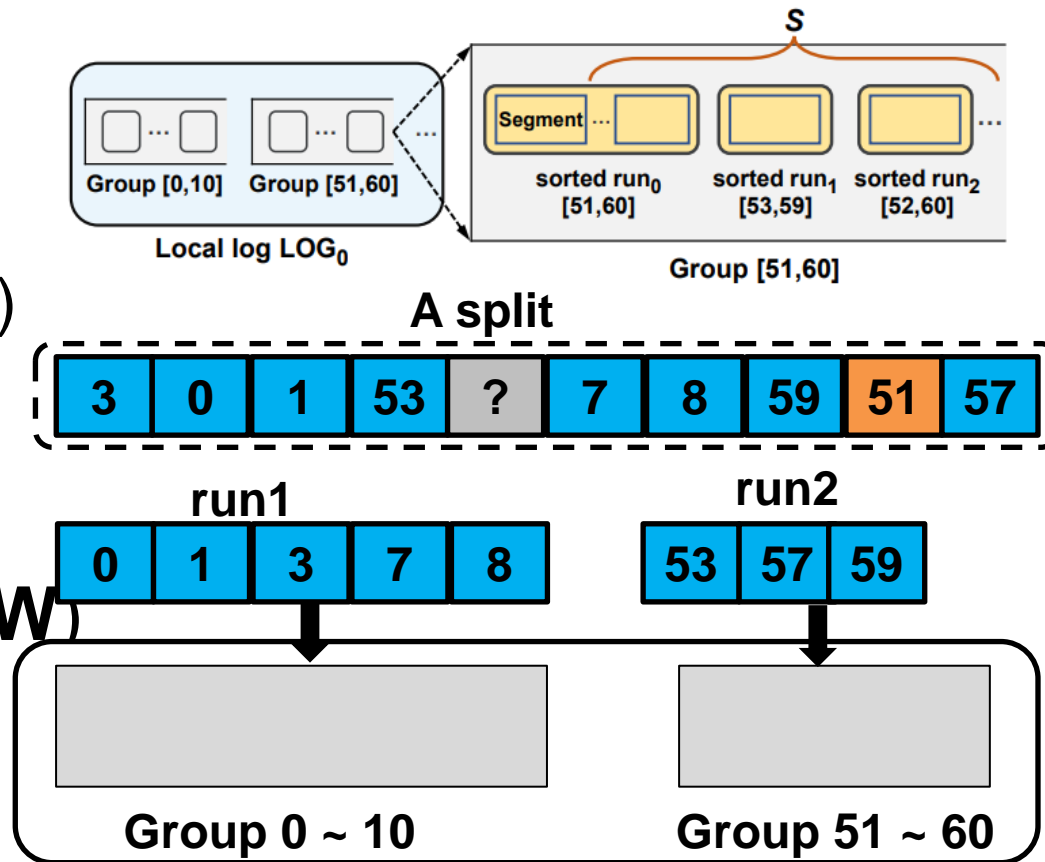
Design

➤ Local Log Structure

- Range-based grouping(virtual nodes-friendly)
- Sorted inside each run
- Each split generate 1 run for each group

➤ Tunable Ordering(balance between R/W)

- user-configurable S: maximum runs in a range group
- Merge new runs into old ones when the S reaches the limit(big overhead)
- S = 1 Read  Write  (Like a simple LSM-Tree, for high RCL)
- S = Inf Read  Write  (for high WCL)



Evaluation

➤ Setup

- Distributed System: local cluster, 10Gb/s switch
- Hardware: 500G SSD 12-core E5-2650 v4

➤ Workloads

- Generated by YCSB(Zipfian constant 0.99)
- Pair size 1KB(key size: 24B)

➤ Settings

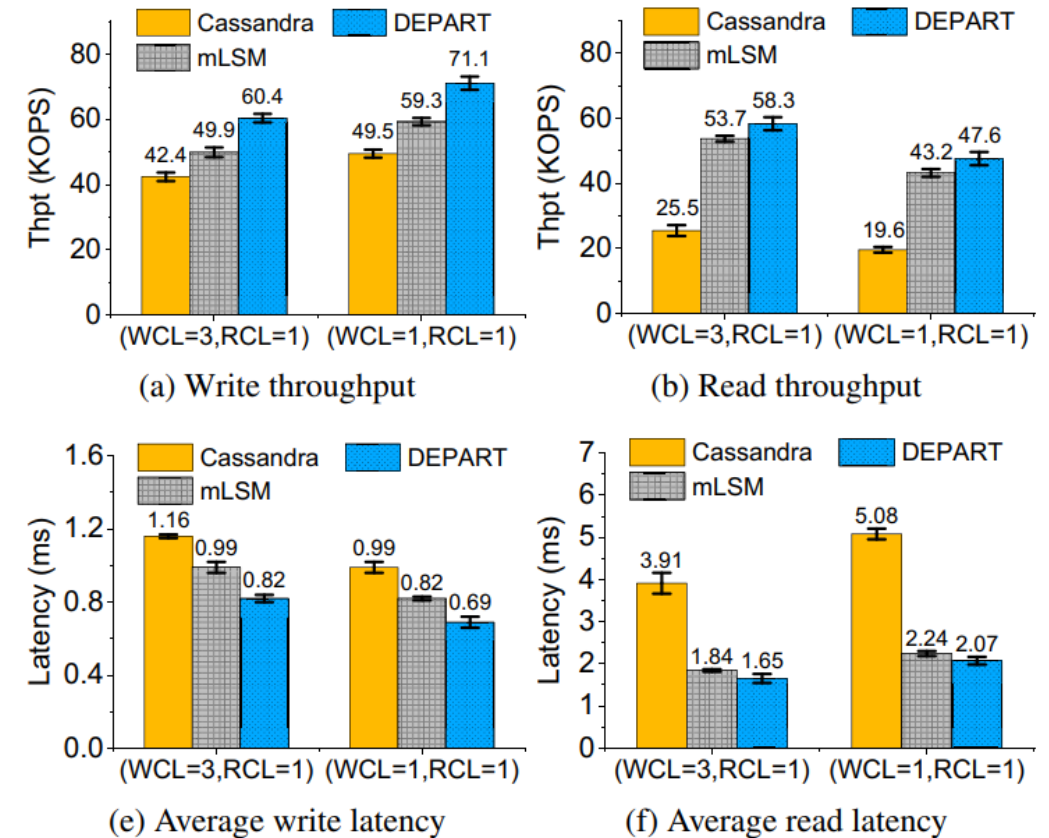
- Compare objects: Cassandra(traditional), mLSM, DEPART
- Storage system: 5 nodes + triple replication
- Different consistency level (WCL= 1,RCL=1 & WCL = 3,RCL=1)

Evaluation

➤ IO Throughputs / Latency

DEPART > mLSM > Cassandra

- DEPART vs. Cassandra
 - Smaller LSM Tree
- DEPART vs. mLSM
 - Smaller gap
 - (2-layer-log > LSM?)(unknown size of runs)
- WCL = 3 vs. WCL = 1
 - (Why does WCL=3 have better read performance)

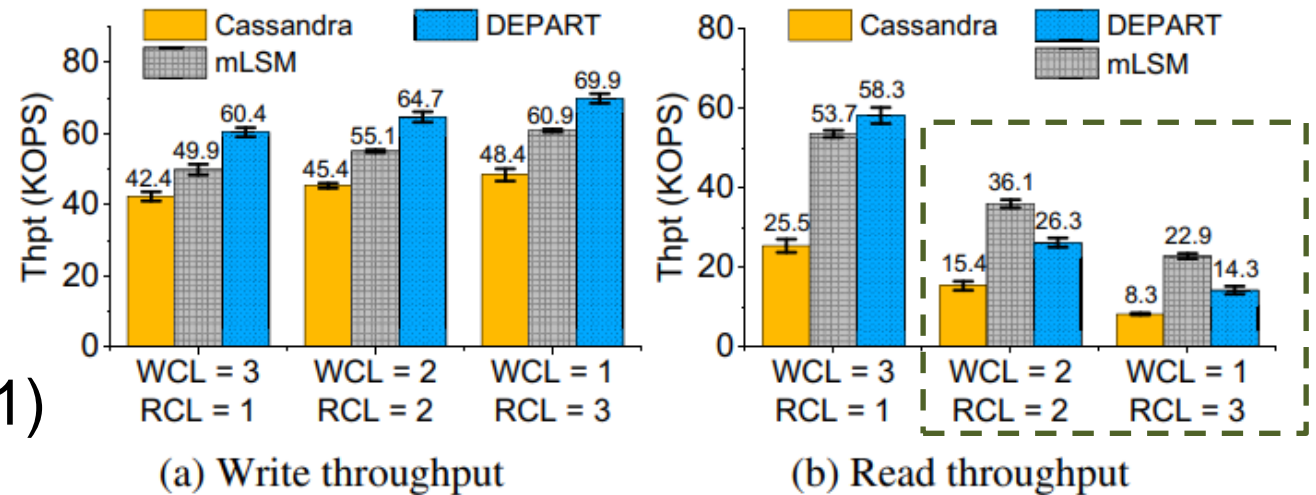


Evaluation

➤ Performance under different consistency configurations

DEPART ? mLSM > Cassandra

- Write
 - DEPART > mLSM
- Read
 - DEPART < mLSM (when RCL > 1)



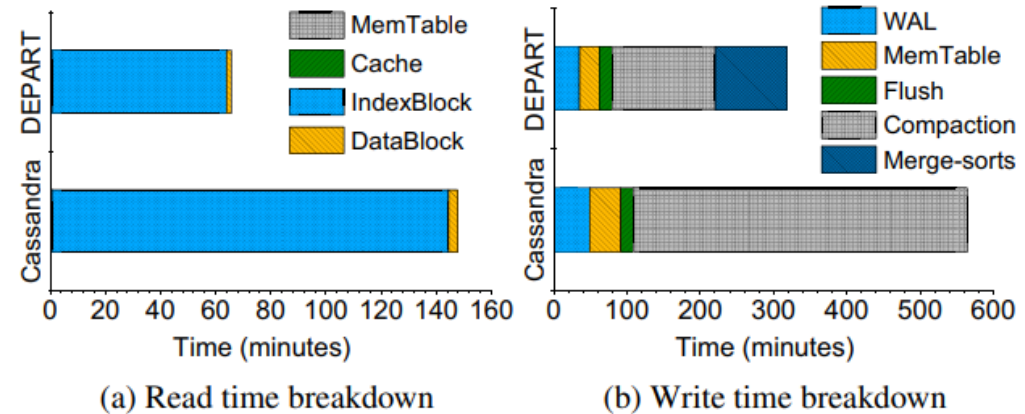
DEPART prefer write performance, It is possible to configure a larger number of runs

Evaluation

➤ Time breakdown No data for mLSM

DEPART > Cassandra

- Read
 - Less index time
- Write
 - Merge sort < compaction



➤ Impact of the ordering degree S

S	Write thpt (KOPS)	Read thpt (KOPS)
1	37.2	42.3
10	57.2	31.5
20	64.7	23.1
$\rightarrow \infty$	78.4	7.6
Cassandra	45.4	15.4

Conclusion

➤ Problem

- Big LSM-Tree cause serious read and write amplification

➤ Idea

- Separation primary data and redundant data
- Trick on consistency level : Balance on Read and Write

➤ Solution

- Store primary data in LSM-Tree
- Store redundant data in 2-layer-log
- 2-layer-log can adjust internal data organization to meet different consistency requirements