

DUPEFS: Leaking Data Over the Network With Filesystem Deduplication Side Channels

Andrei Bacs and Saidgani Musaev, VUSec, Vrije Universiteit Amsterdam;

Kaveh Razavi, ETH Zurich; Cristiano Giuffrida and Herbert Bos,

VUSec, Vrije Universiteit Amsterdam

FAST 2022

Background

➤ Filesystem&Dedup

- Inline Deduplication
- Basic write workflow
- Data identifiers
- Deduplication tables
- Deduplication granularity

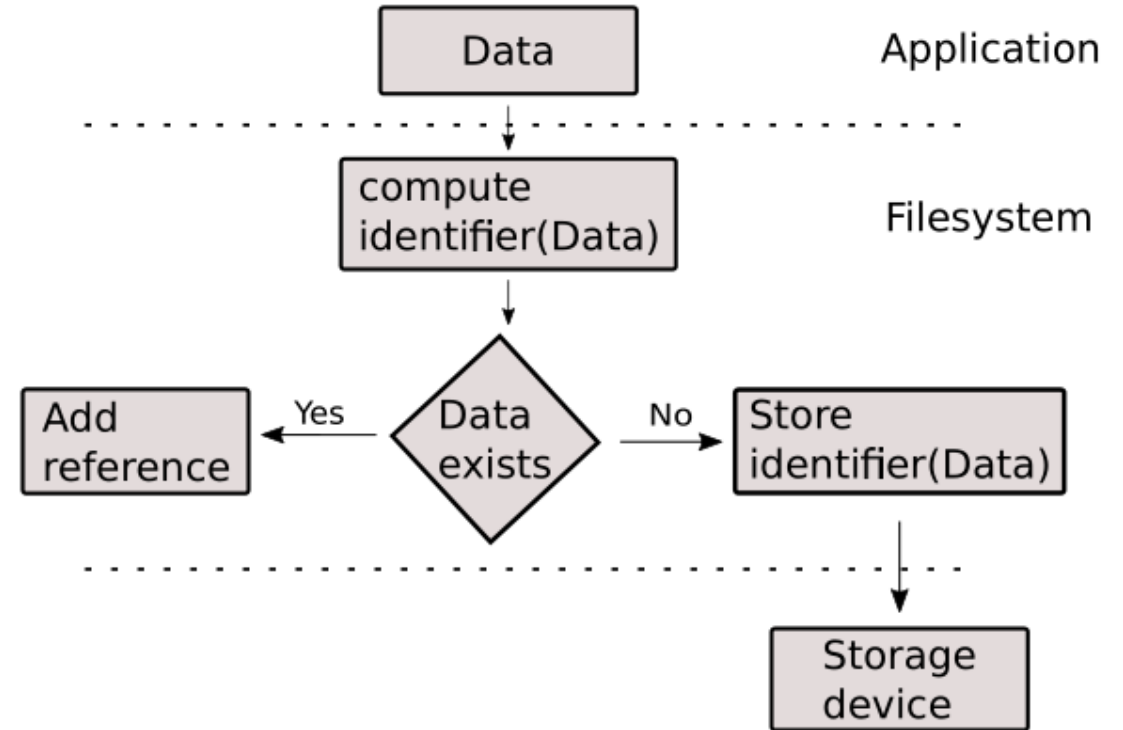

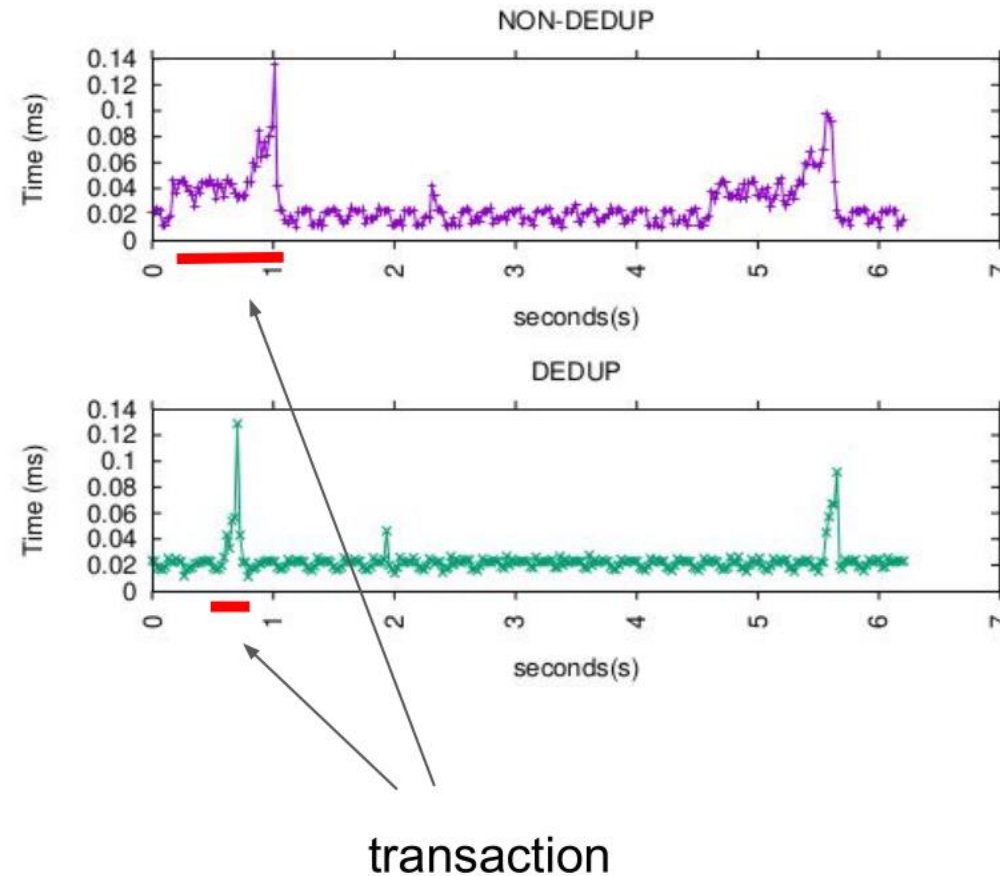


Figure 1: Write path with deduplication.

Deduplication timing side channel



Write path with deduplication



Transaction is shorter (pulse width) for duplicate data

Background

➤ ZFS

- Transactional copy-on-write filesystem
- Deduplication record(128KB)&deduplication table(DDT)

➤ Btrfs

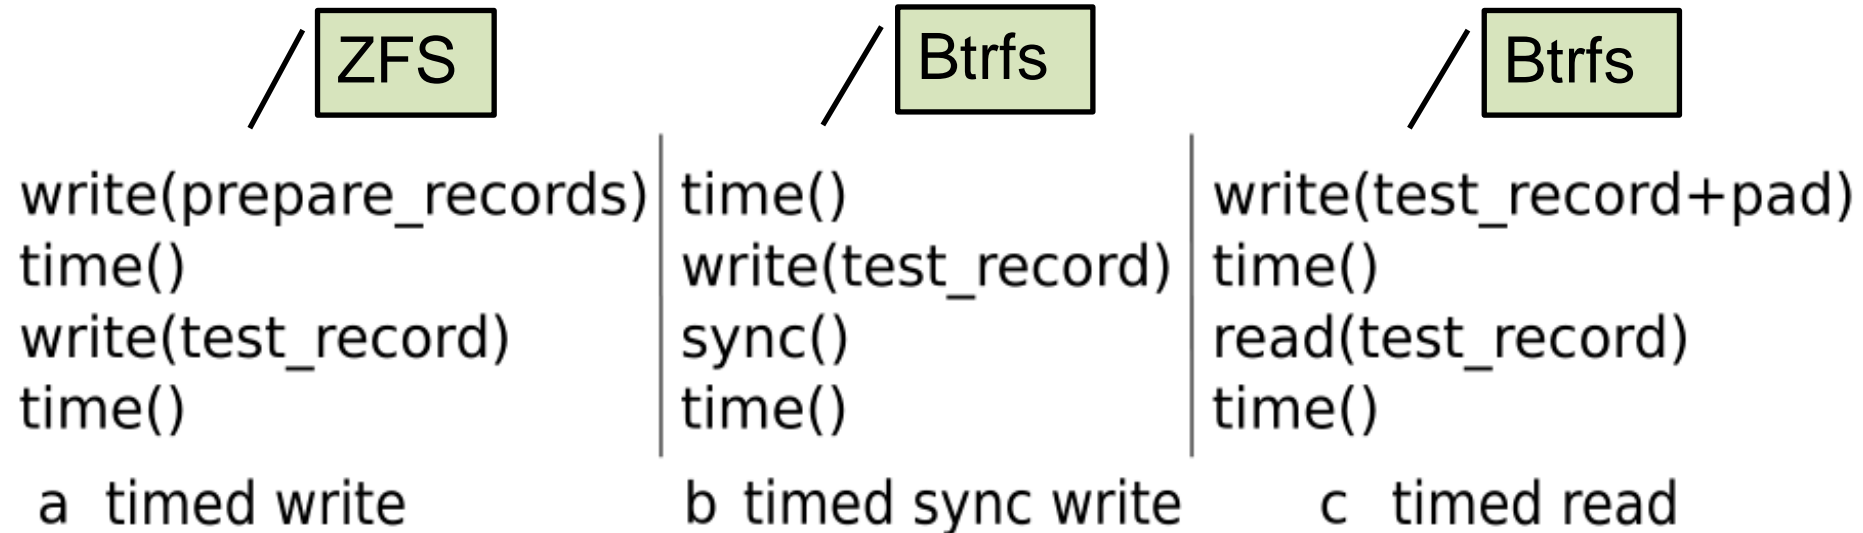
- Extent(128KB): The data are stored in extents
- COW filesystem

➤ Primitive

- The timed write primitive
- The timed read primitive

Attack Primitives

- carefully-chosen sequence of file operations



Attack Challenges

➤ **Filesystem asynchronous I/O operations**

- intermediary caches
- transactional behavior
- ✓ Exploitation technique: filesystem cache massaging

➤ **Deduplication granularity**

- typical record size 128KB (ZFS and Btrfs)
- ✓ Exploitation technique: alignment probing

➤ **Weak amplification in a remote attack**

- ✓ Exploitation technique: secret spraying

Threat Model

- **Attacker and victim have access to the same filesystem**
- **The filesystem uses inline deduplication and default settings**
- **No limit on I/O operations**

DUPEFS Exploitation

➤ Data fingerprinting

- reveal the presence of existing known but inaccessible data

➤ Data exfiltration

- exfiltrate secret data from a system (or sandbox)

➤ Data leak

- leak secret data from a remote system
- alignment probing
- secret spraying

Alignment probing

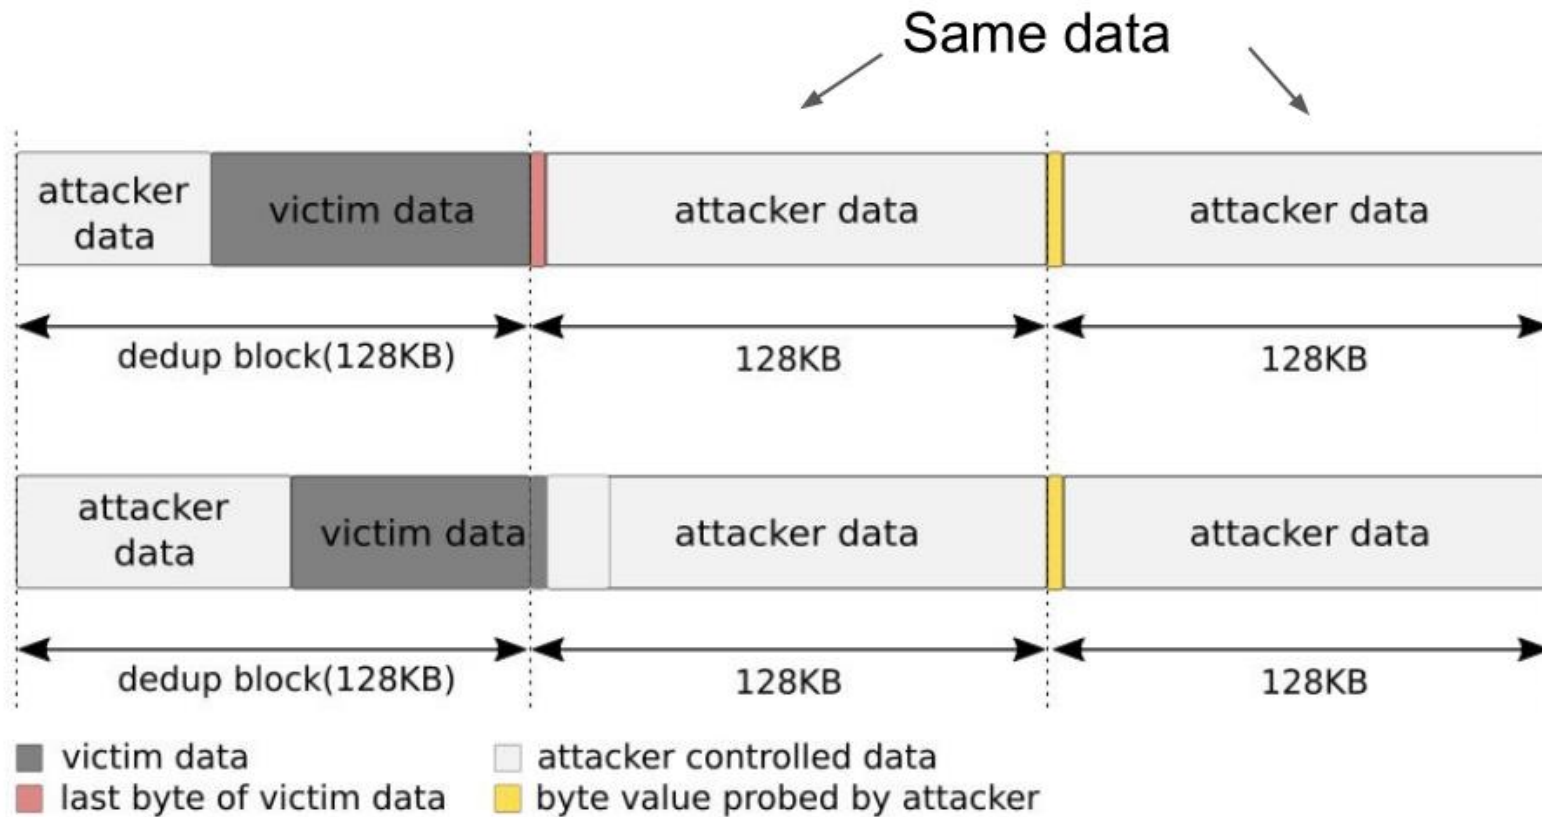
- Enables byte granularity
- Reduces entropy

Timed write primitive

```
write(prepare_records)  
time( )  
write(test_record)  
time( )
```




Alignment probing



Secret spraying

- Amplifies the timing signal
- N dedup events per correct guessed byte




Data Fingerprinting/exfiltration

➤ Data Fingerprinting

- Btrfs-based synchronous write primitive
- local exploitation scenario

➤ Data exfiltration

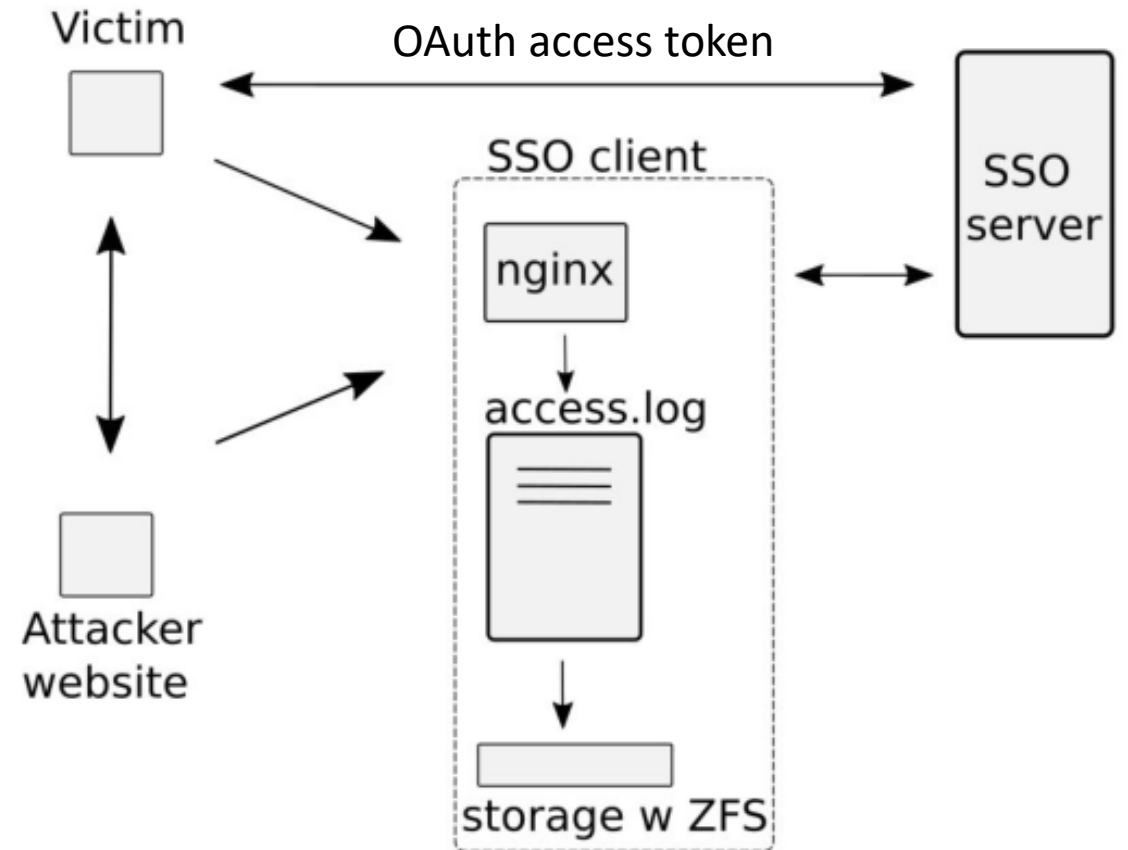
- Btrfs-based synchronous write primitive
- A local unprivileged attacker (or “sender”) seeks to exfiltrate data from a sandbox over a covert channel.



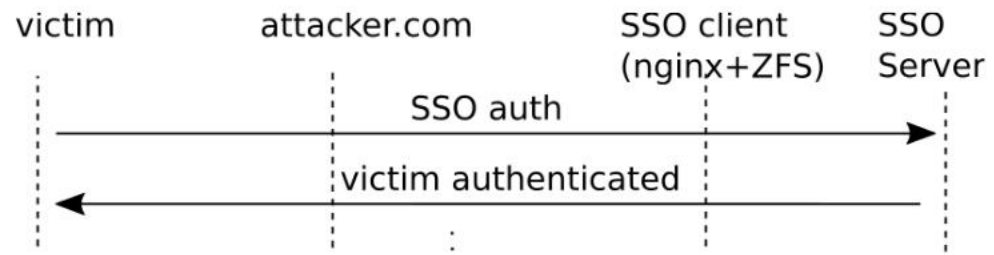
Local: data fingerprinting, exfiltration

Data leak

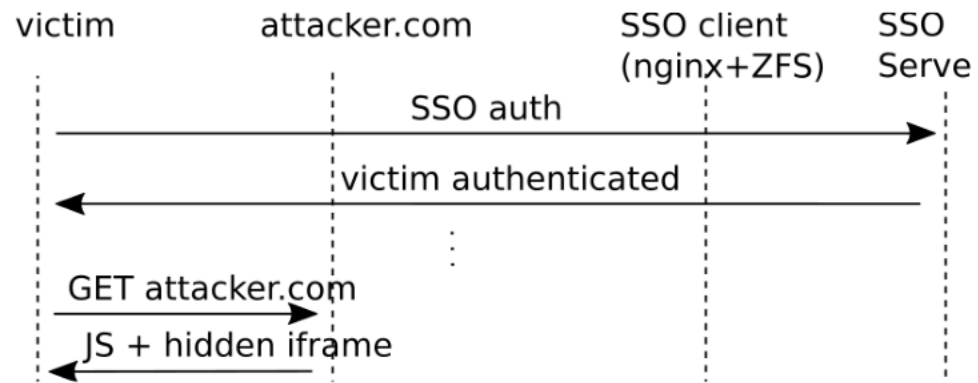
- targets access tokens of OAuth 2.0 implicit grant access scheme
- SSO client runs nginx on top of ZFS and logs requests
- access tokens are encoded in the request and do not expire during attack
- SSO client does not offer X-Frame-Options



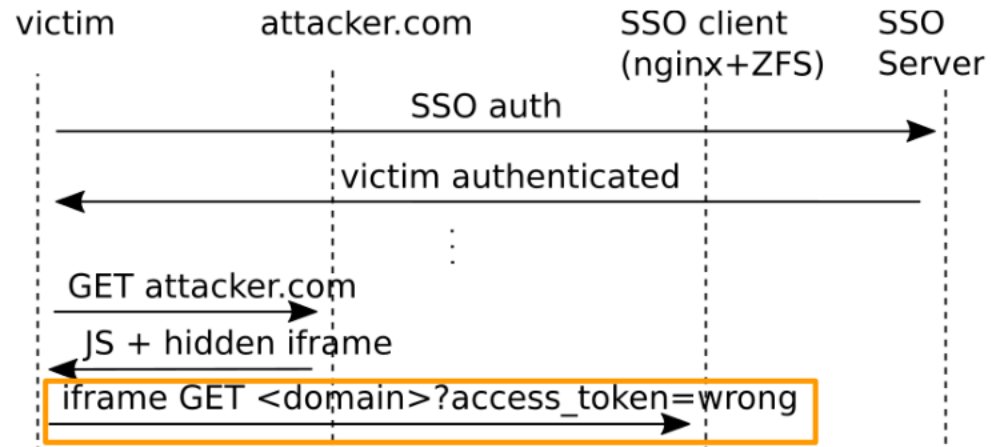
Data leak




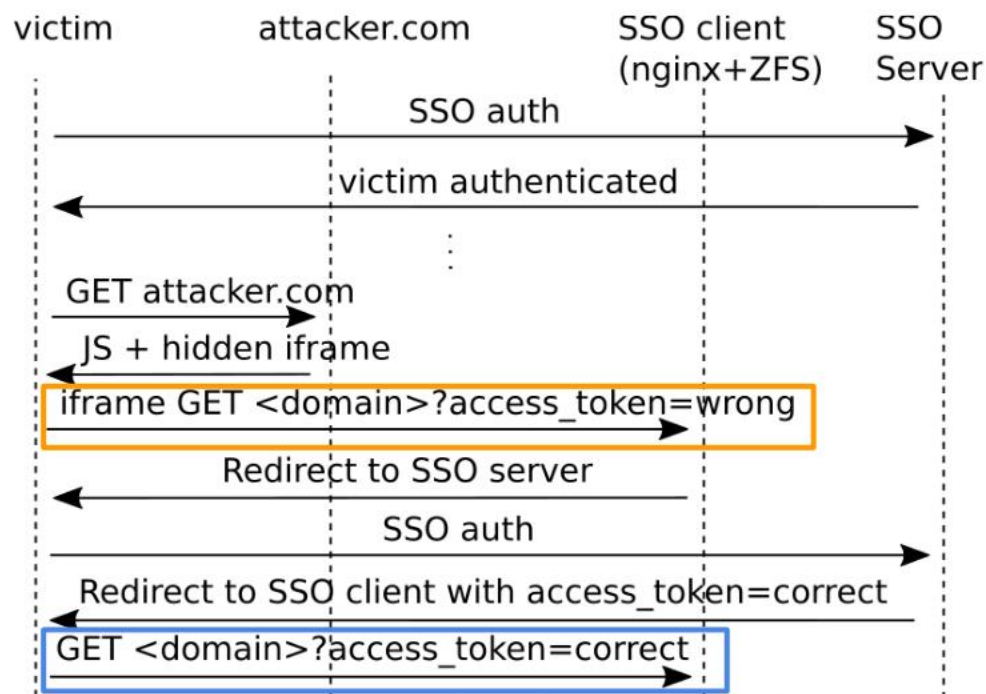
Data leak



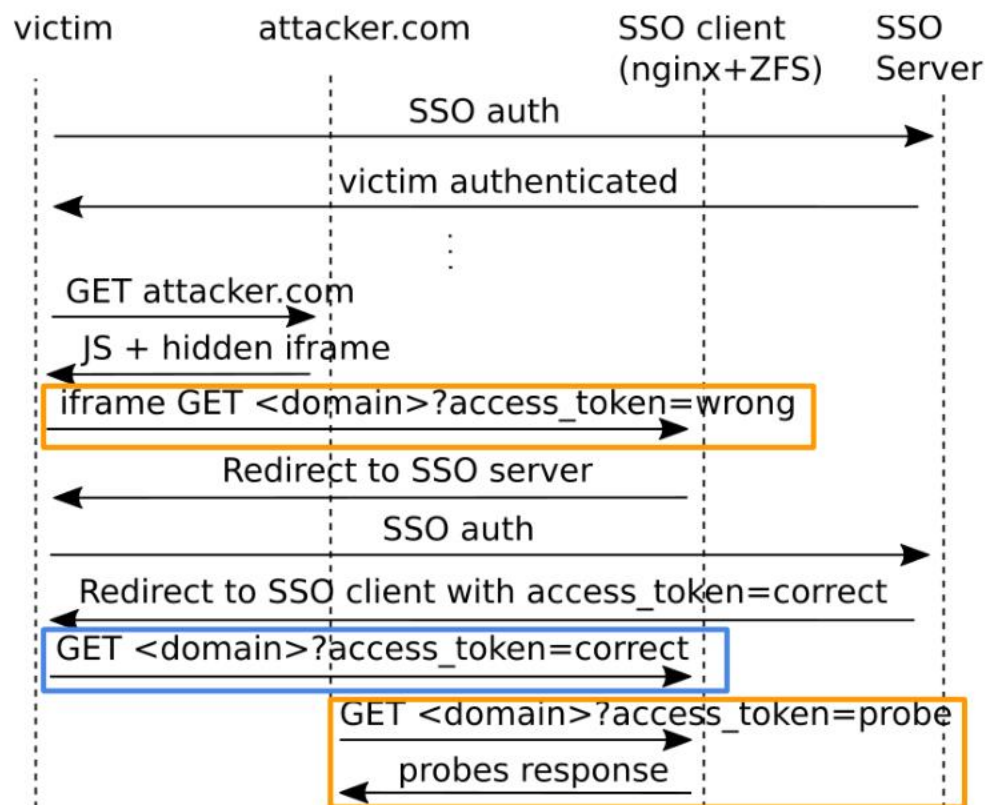
Data leak



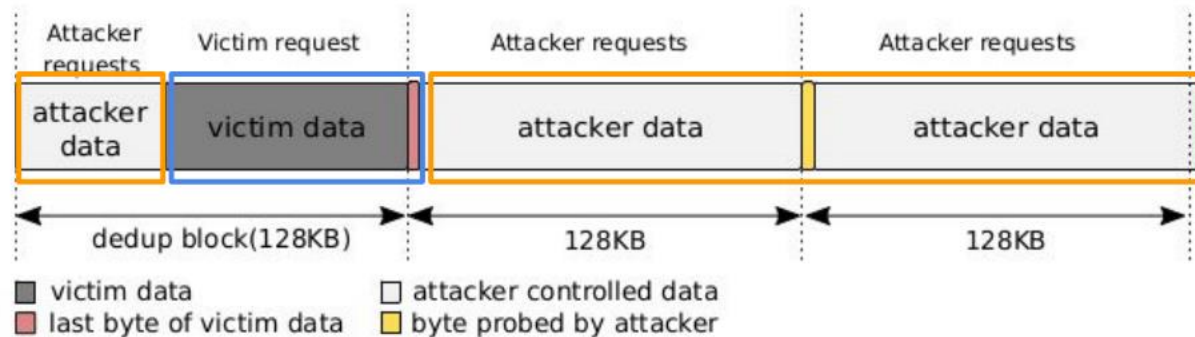
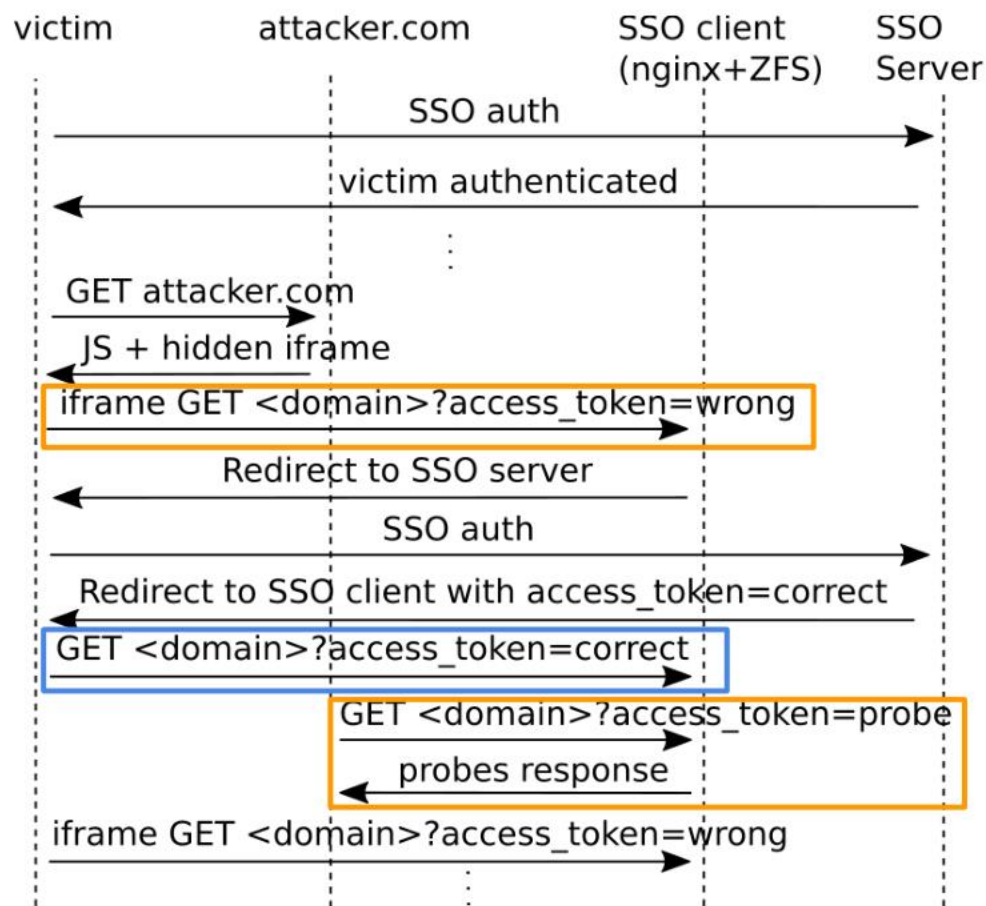
Data leak



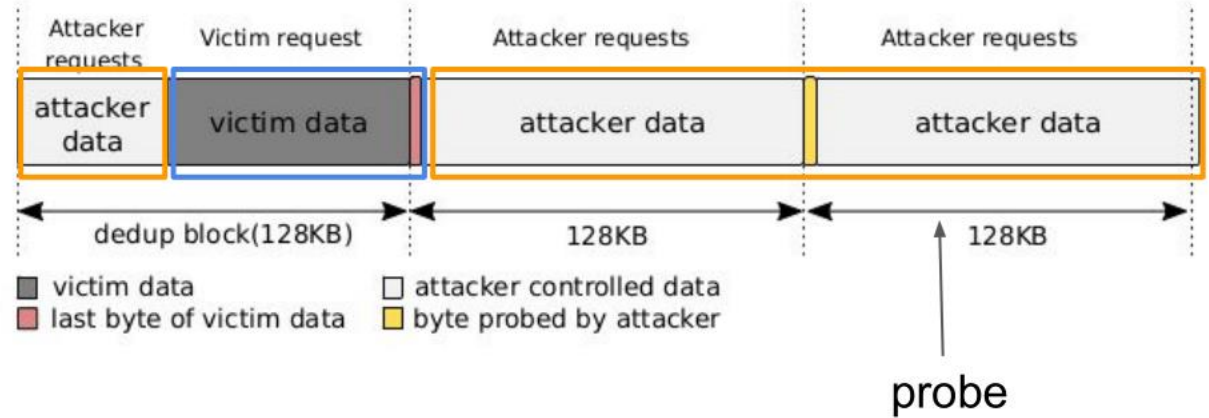
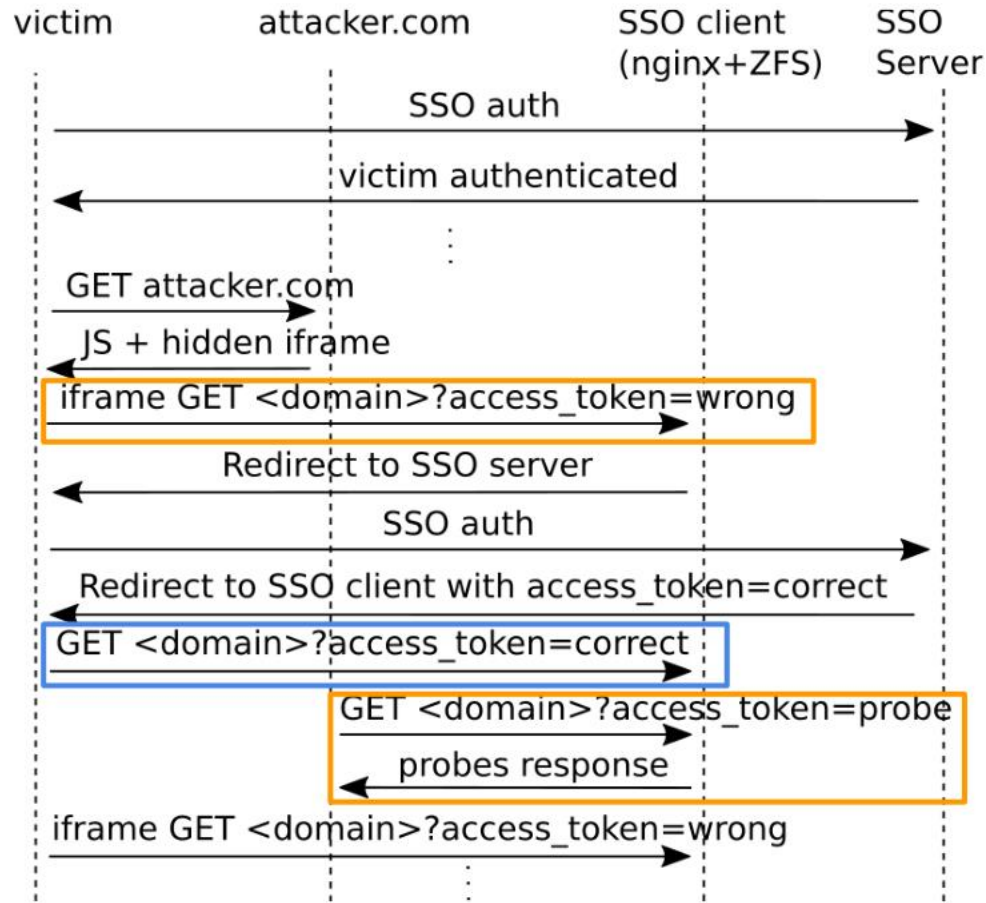
Data leak



Data leak



Data leak



- correct byte value probe produces transaction durations < threshold

Evaluation

➤ Data fingerprinting

- DUPEFS can reliably fingerprint the target data except the last sub-128 KB chunk of a file. Thus the small (181 KB and 223 KB) files have lower success rates.

Table 1: File fingerprinting

| File | Type | Size | Success |
|-------------------------------|--------|--------|---------|
| config-4.11.3-200.fc25.x86_64 | text | 181 KB | 70% |
| lena_color.gif | binary | 223 KB | 55% |
| libz3.so | binary | 22 MB | 99% |
| x86_64-redhat-linux-c++ | binary | 1 MB | 99% |

Evaluation

➤ Data Exfiltration

Table 2: Covert channel

| N | Bit errors | Time | BR | BER | I/O |
|-----|------------|-------|-----------|--------|---------|
| 20 | 13 | 375s | 0.320 bps | 10.83% | 76.8MB |
| 40 | 14 | 746s | 0.160 bps | 11.66% | 153.6MB |
| 60 | 12 | 1591s | 0.075 bps | 10.00% | 230.4MB |
| 100 | 6 | 1873s | 0.064 bps | 5.00% | 384.0MB |
| 120 | 3 | 2387s | 0.050 bps | 2.50% | 460.8MB |

Evaluation

➤ Data Leak

- OAuth token 22 bytes (base64)
- LAN: 1 hop (RTT 0.1ms)
- WAN: 12 hops (RTT 2ms)

Table 3: LAN 1 byte data leak

| Success | Attack time/byte | Probes/byte val | I/O |
|---------|------------------|-----------------|---------|
| 50% | 19.2 min | 200 | 4.9 GB |
| 80% | 25.6 min | 300 | 7.3 GB |
| 92% | 42.6 min | 400 | 9.8 GB |
| 96% | 78.9 min | 800 | 19.6 GB |

Table 4: WAN 1 byte data leak

| Success | Attack time/byte | Probes/byte val | I/O |
|---------|------------------|-----------------|---------|
| 64% | 24.5 min | 200 | 4.9 GB |
| 87% | 38.4 min | 300 | 7.3 GB |
| 94% | 59.7 min | 400 | 9.8 GB |
| 94% | 110.9 min | 800 | 19.6 GB |

Mitigation

- **Deduplication ideal implementation**
 - Save space
 - Constant time behavior
- **Pseudo-same behavior policy**
 - Perform data overwrite for duplicate data
 - Read path: time jitter
 - Renders remote attacks impractical

Contributions

- Analyze filesystem deduplication side channels, show that despite the challenges, attackers can mount byte-level data leak attacks across the network.
- Introduce DUPEFS's novel attack primitives and demonstrate their feasibility in end-to-end attacks to leak data even across the Internet.
- Describe and analyze mitigations for such attacks.

Conclusion

