

Building a High-performance Fine-grained Deduplication Framework for Backup Storage with High Deduplication Ratio

Xiangyu Zou and Wen Xia, Harbin Institute of Technology, Shenzhen; Philip Shilane, Dell Technologies; Haijun Zhang and Xuan Wang, Harbin Institute of Technology, Shenzhen

Speaker wrl

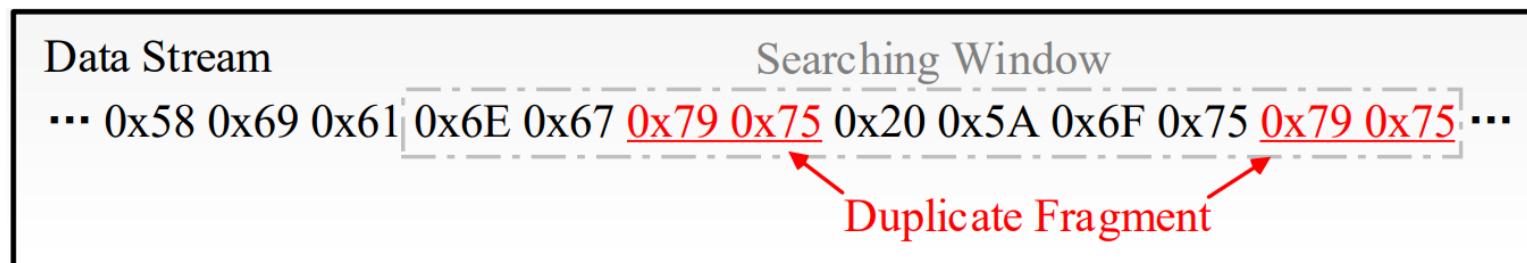
ATC 2022

Background

- **Data Reduction:** the purpose is to reduce the physical capacity required to store the growing amounts of logical backup data.

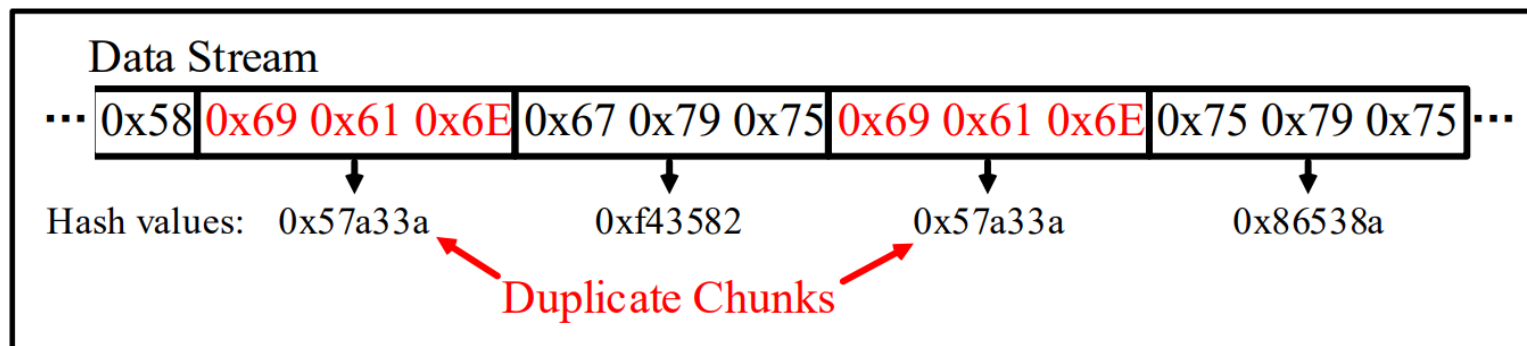
- **General Compression :**

- For usual-size files
- String-level
- Limited window



- **Deduplication:**

- For very large files
- Chunk-level
- Global

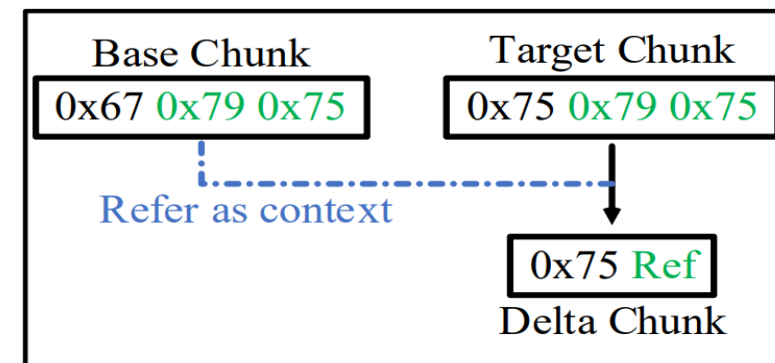
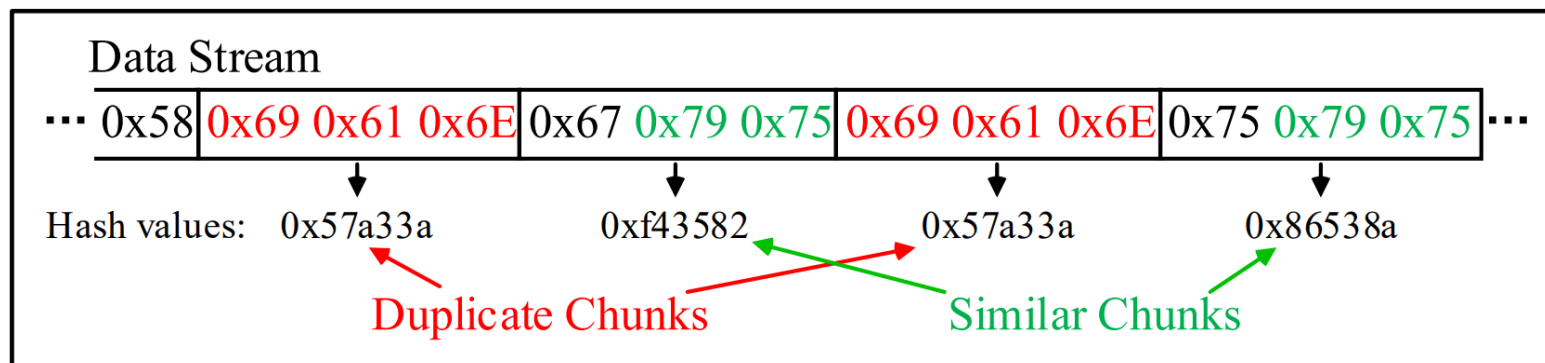


- Both have been widely used in storage products
- Can not fully utilize data's compressibility

Background

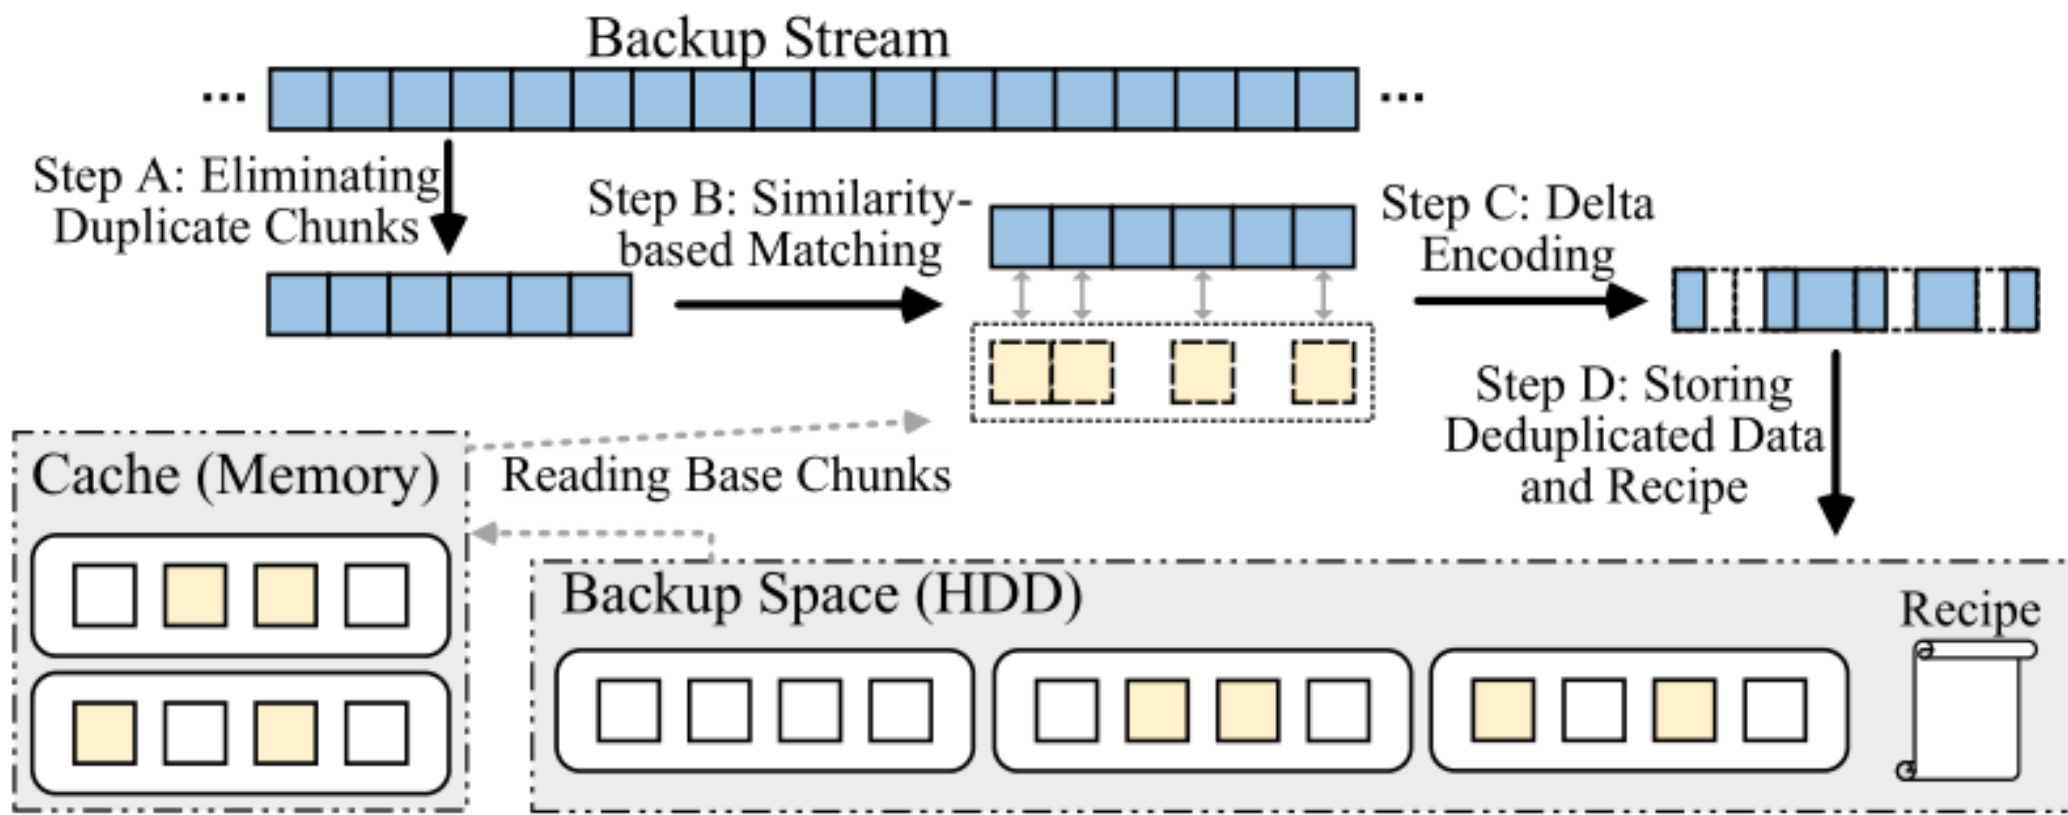
➤ Fine-grained deduplication:

- Leverages not only identical chunks, but also similar chunks
- Introduces additional steps on post-deduplication chunks
 - Detects similar chunks for an unduplicated chunk (i.e., target chunk for delta encoding)
 - Reads back a similar one (an already stored chunk) as a base chunk
 - Calculates delta difference between the target chunk and the base chunk
- String-level, Global



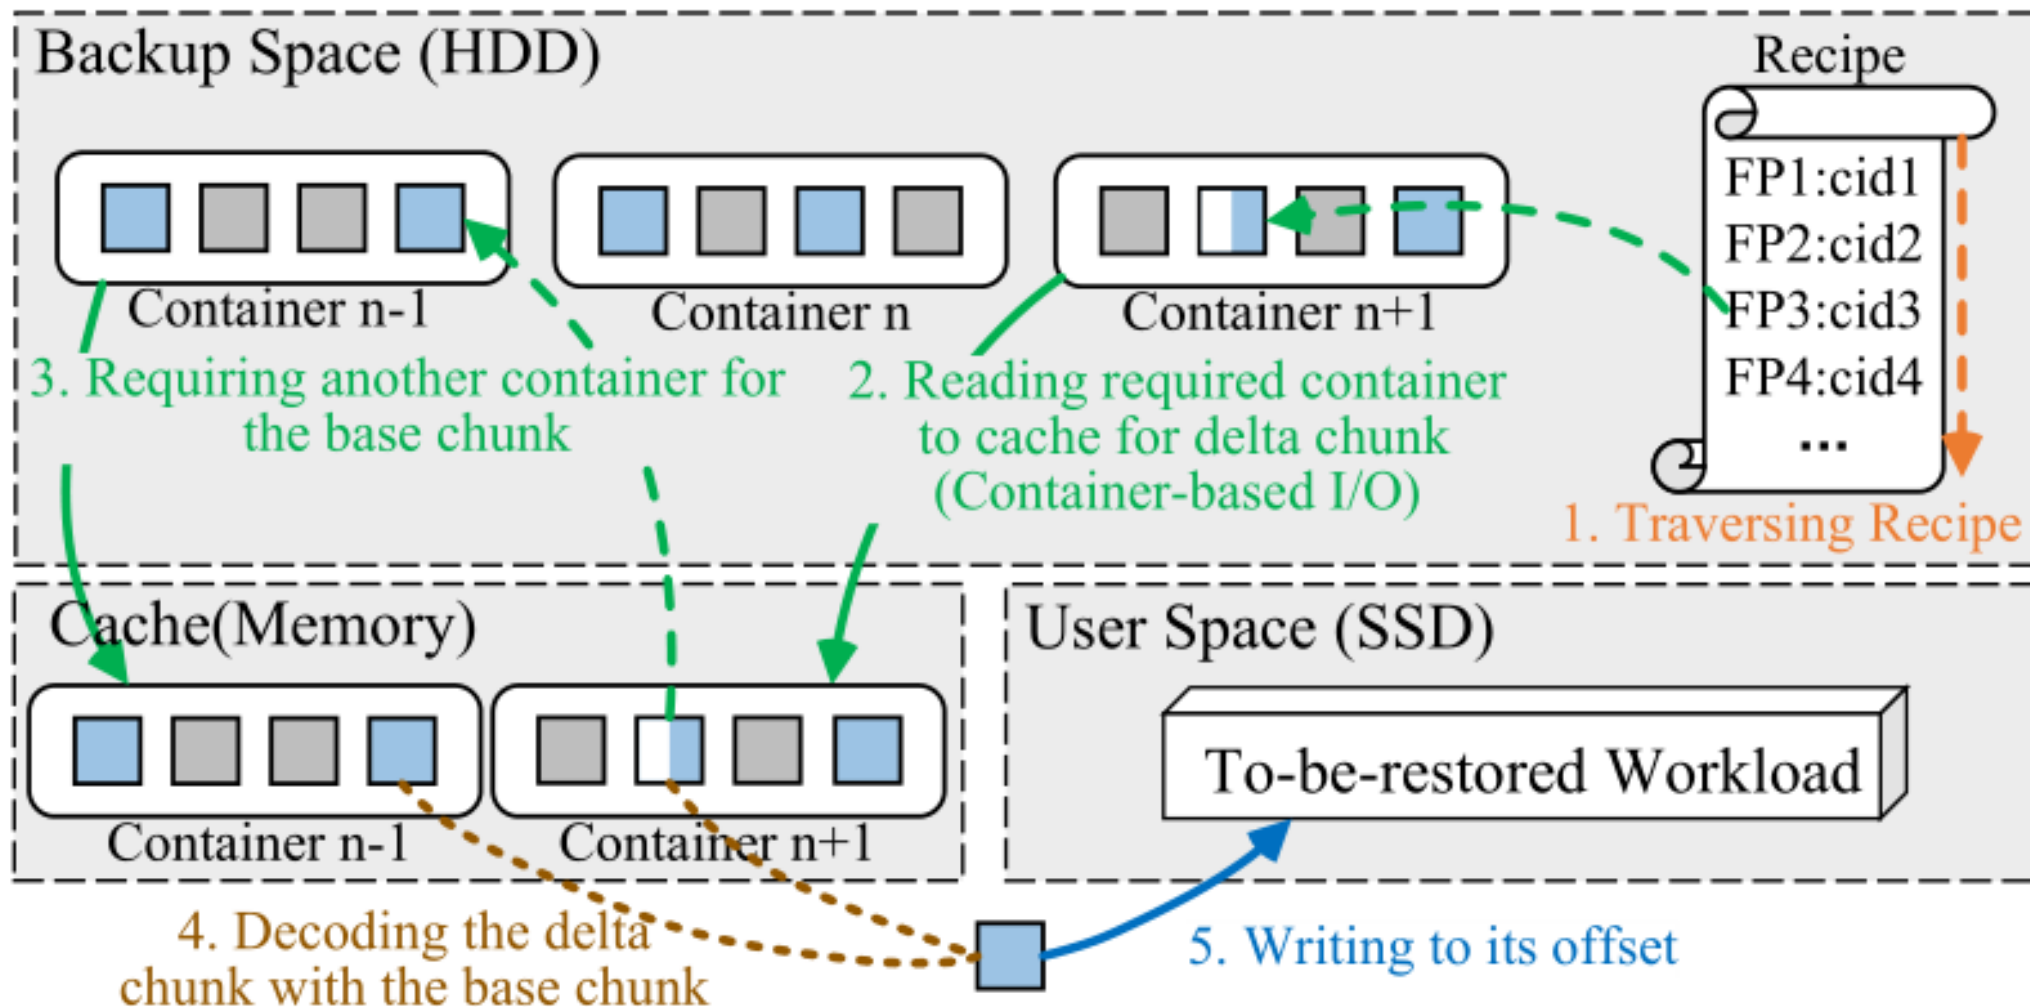
Background

- The backup workflow of fine-grained deduplication.



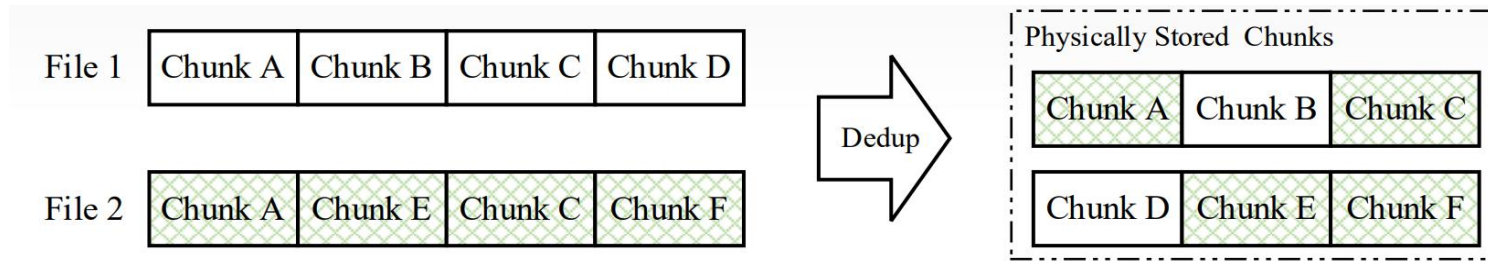
Background

- Restoring a delta chunk in the restore workflow of fine-grained deduplication.



Problem

- **Dedupcation: Reusing data hurts locality ,declines systems' performance.**



- **The problem:** Fine-grained deduplication introduces a new form of data reuse, make some **additional locality issues**

- **Challenge 1:** Poor locality in the backup workflow causes inefficient I/O when reading base chunks. —**Poor locality in base chunks (the write path)**
- **Challenge 2:** Delta-base relationships lead to more complex fragmentation problems than deduplication alone.—**Poor locality in restore-required chunks (the read path)**
- **Challenge 3:** Delta-base dependencies cause poor temporal locality during delta decoding and causes repeated container reads. —**Poor locality in delta-base pairs (the read path)**

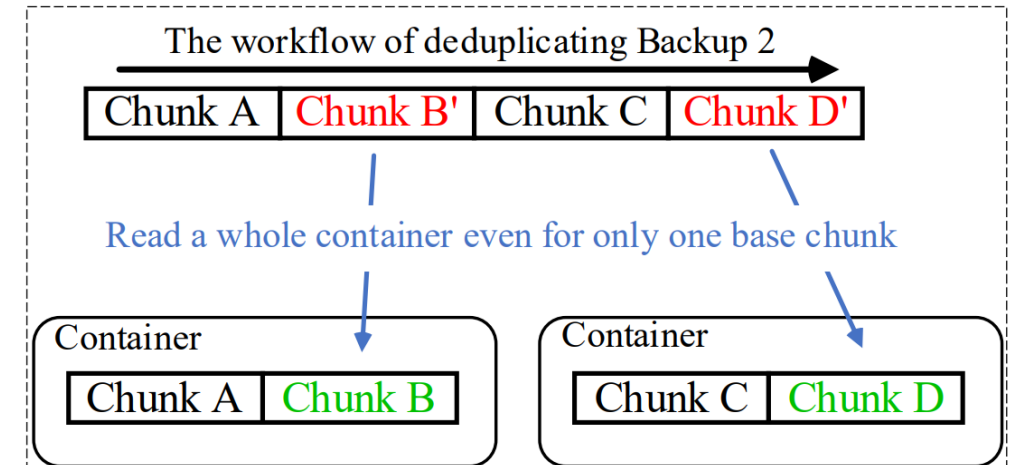
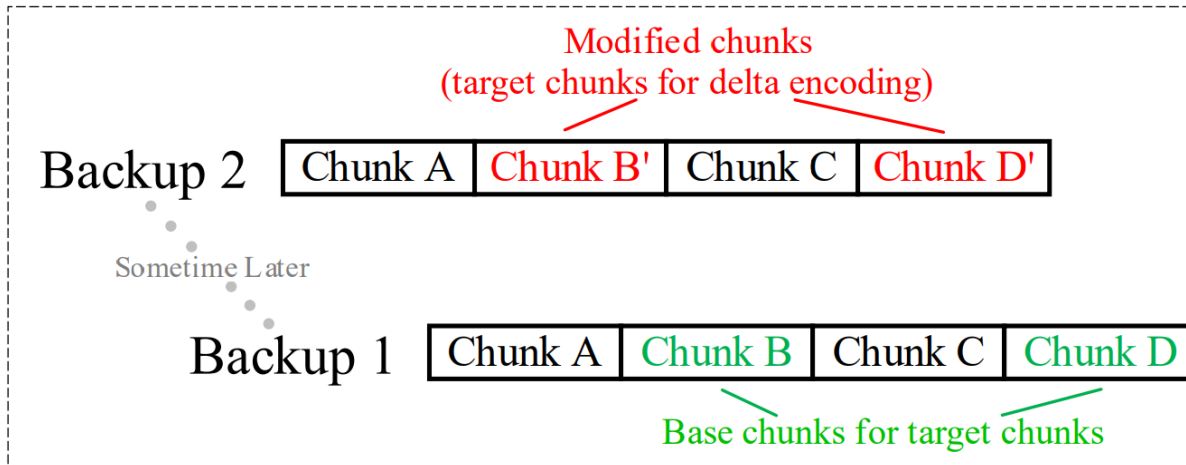

Challenge Poor locality in base chunks (the write path)

➤ Causes:

- The distribution of base chunks' physical positions is random
- Consecutive chunks are usually compressed together (local compression)
 - Accessing the whole compression unit (e.g., container) even for only one chunk

➤ Results:

- Need to read a whole container even for only one base chunk
- Inefficient I/O when reading base chunks in the write path




Challenge Poor locality in restore-required chunks (the read path)

➤ Causes:

- Two kinds of reference relationships
 - Backup workloads – Chunks (introduced by chunk-level deduplication)
 - Delta chunks – Base chunks (additionally introduced by delta encoding)
 - Aggravate the fragmentation problem
- Local compression leads to a large I/O unit

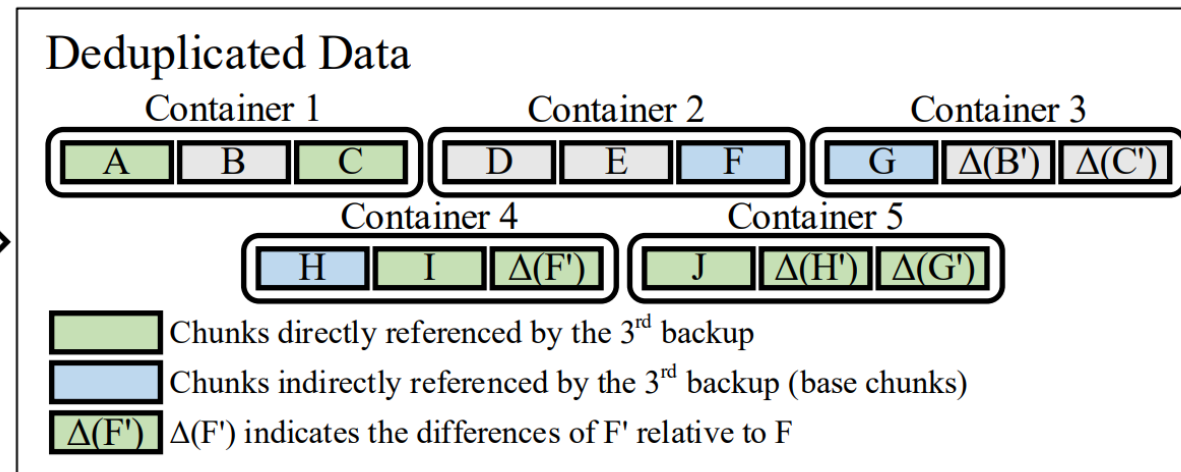
➤ Results:

- Inefficient I/O when reading restore-required chunks in the read path




C1 C3 C1 C3 C1 C4 C2 C3

Fine-grained
Deduplication



Challenge

Poor locality in delta-base pairs (the read path)

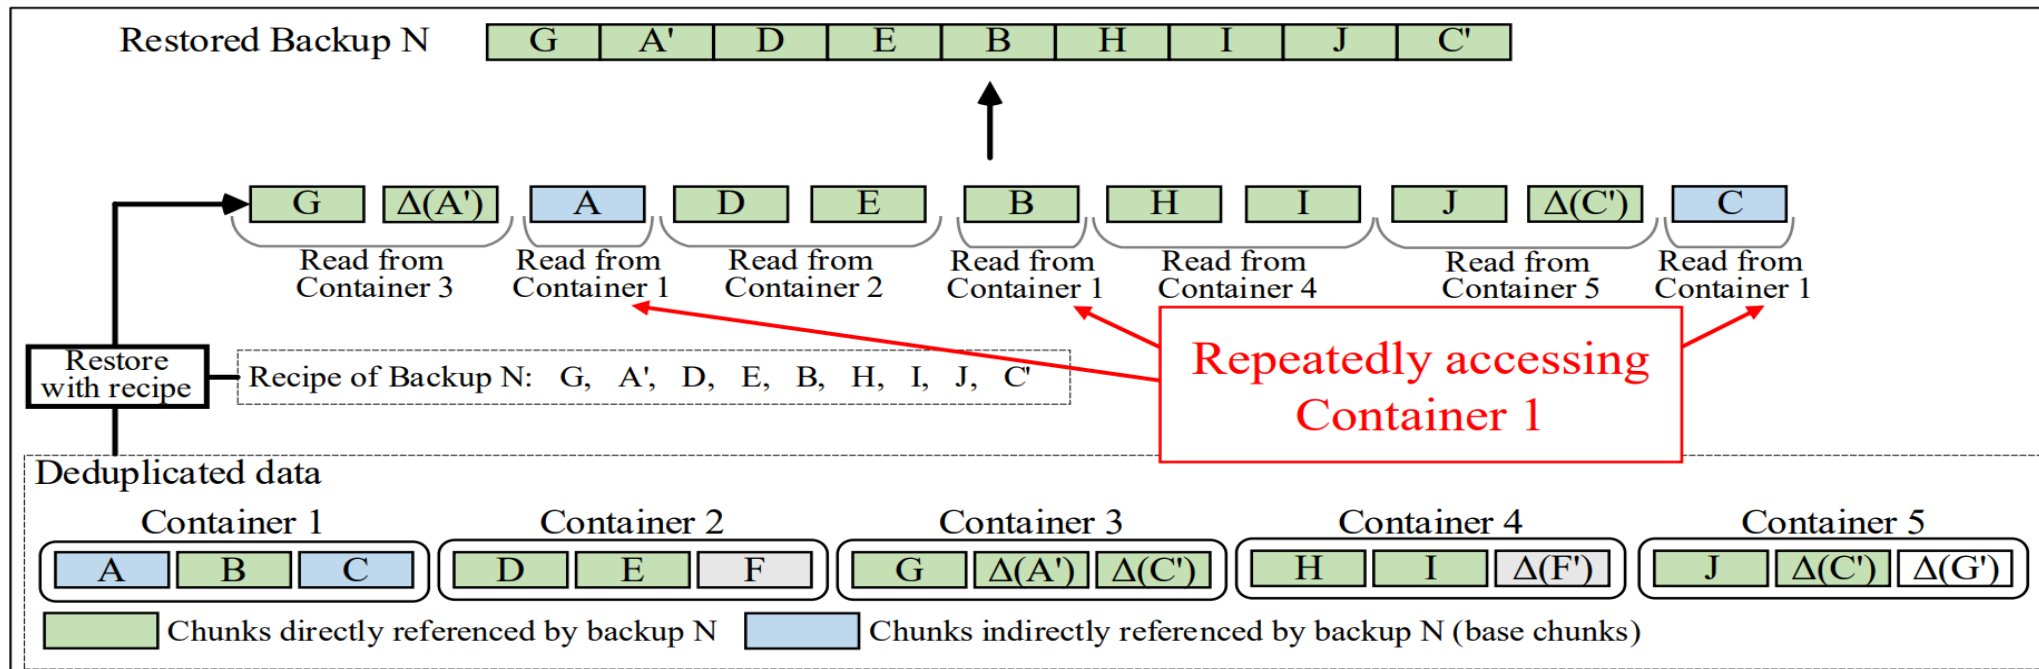


➤ Causes:

- Traversing restore-required chunks when restoring a deduplicated backup
- Delta chunks have dependencies, but usually are far away from their bases

➤ Results:

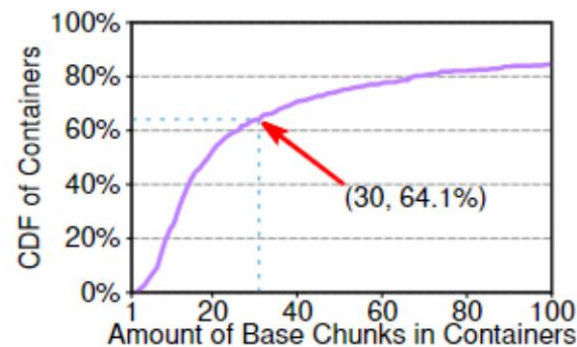
- Repeatedly accessing containers in the read path




Design Selective Delta Encoding

➤ Key Idea: Skip delta encoding if base chunks are in base-sparse containers

- An observation: Base chunks are not distributed evenly
- For example, in an evaluated dataset:
 - 64.1% containers hold ~30 base chunks (“base-sparse containers”)
 - These 64.1% containers only includes 8.31% of the total base chunks.
- Avoids reading these “inefficient” containers in the deduplication workflow



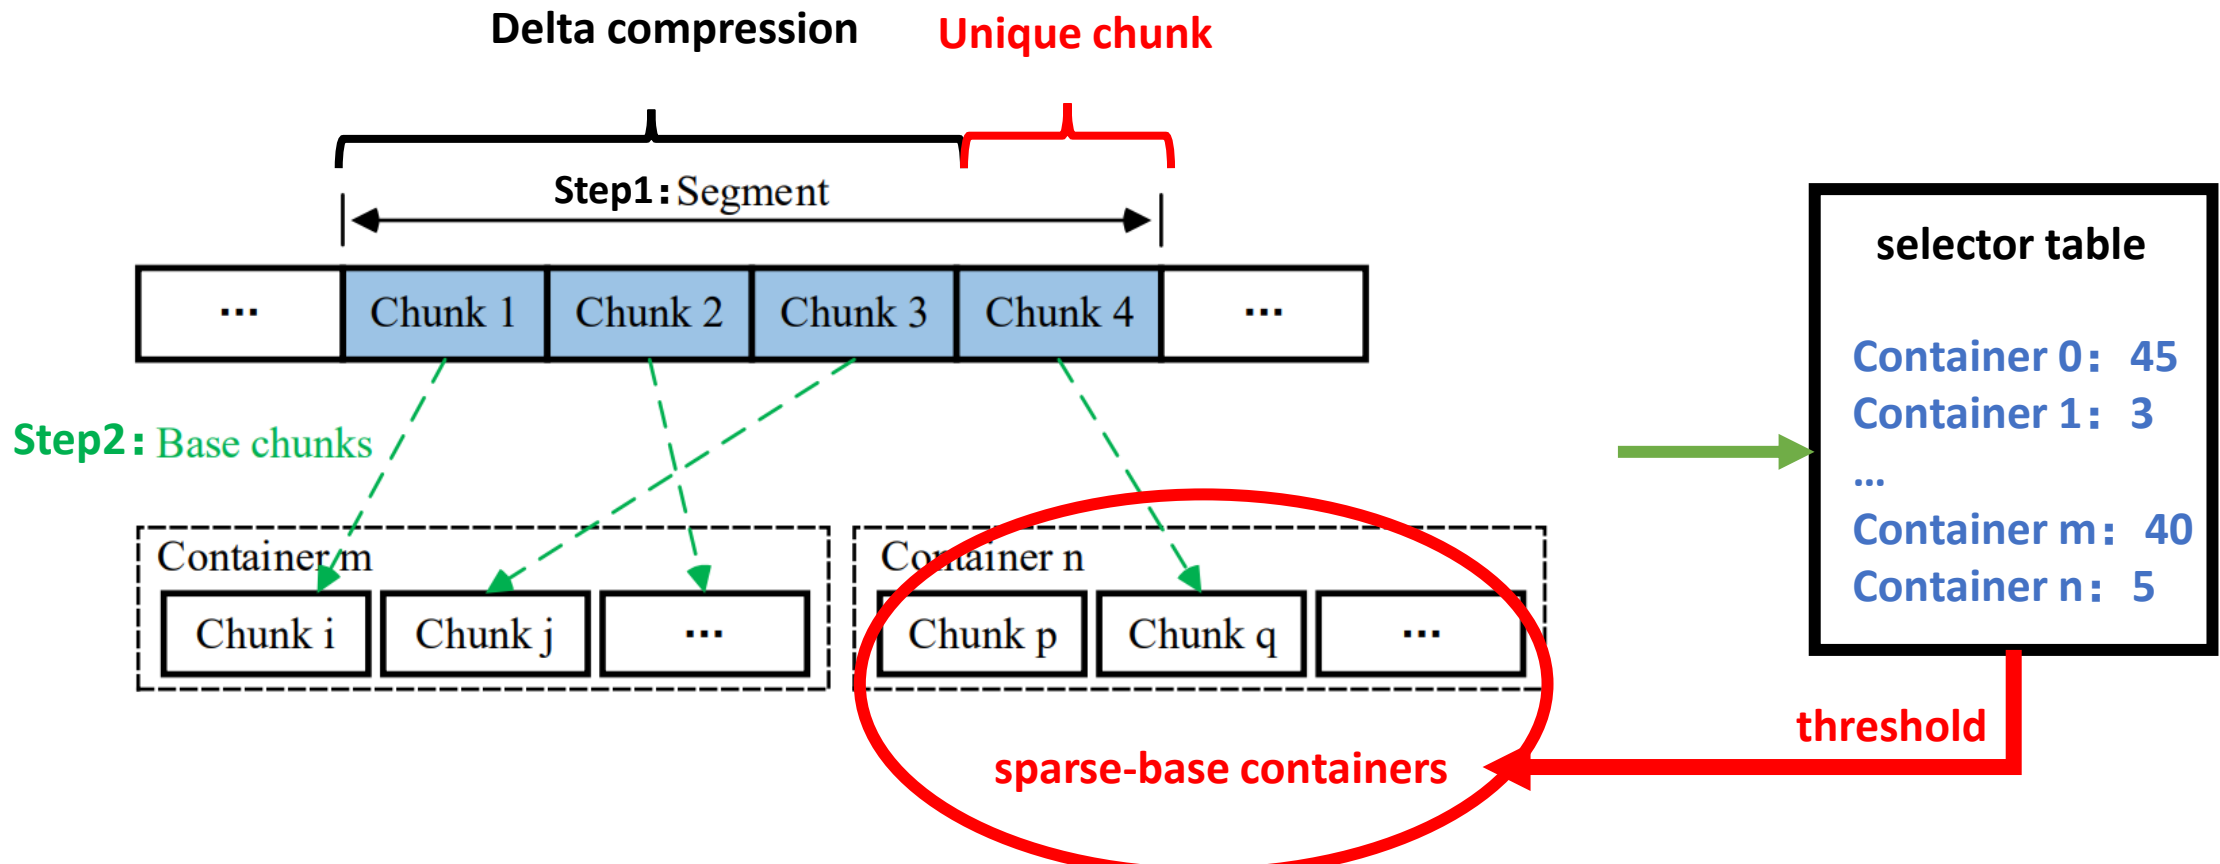
(a) 64.1% of containers contain only ~30 base chunks.



(b) These 64.1% containers only includes 8.31% of the total base chunks.

Design Selective Delta Encoding

- **Key Idea: Skip delta encoding if base chunks are in base-sparse containers → lower I/O and lower compress ratio.**

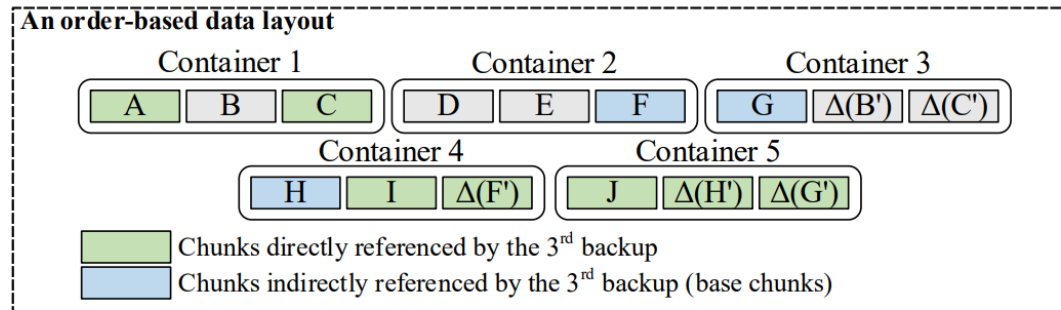


Design Delta-friendly Data Layout

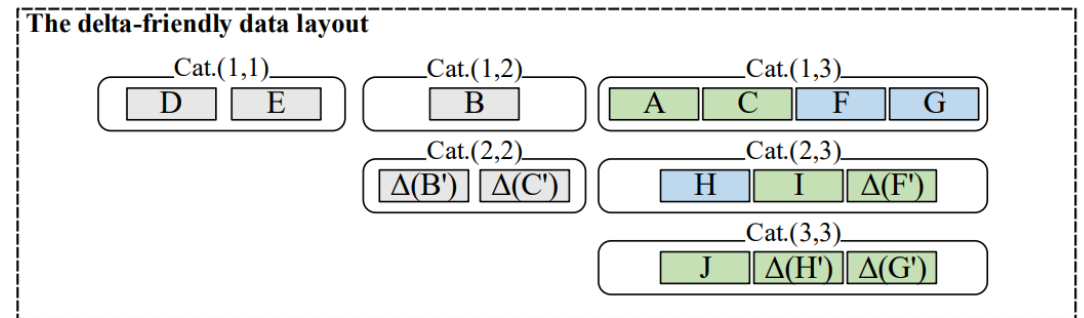
➤ Key Idea: Put the chunks from the same backup together

- Consider two kinds of reference relationship
 - The “Necessary Chunks” of a backup
 - The combination of a backup’s directly and indirectly referenced chunks
- The lifecycle of a chunk
 - A set of backups whose "Necessary chunks" includes this chunks.
- Lifecycle-based classification
- Avoids reading sparse containers in the restore workflow

The 1 st backup	A	B	C	D	E	F	G
The 2 nd backup	A	B'	C'	H	I	F'	G
The 3 rd backup	A	J	C	H'	I	F'	G'

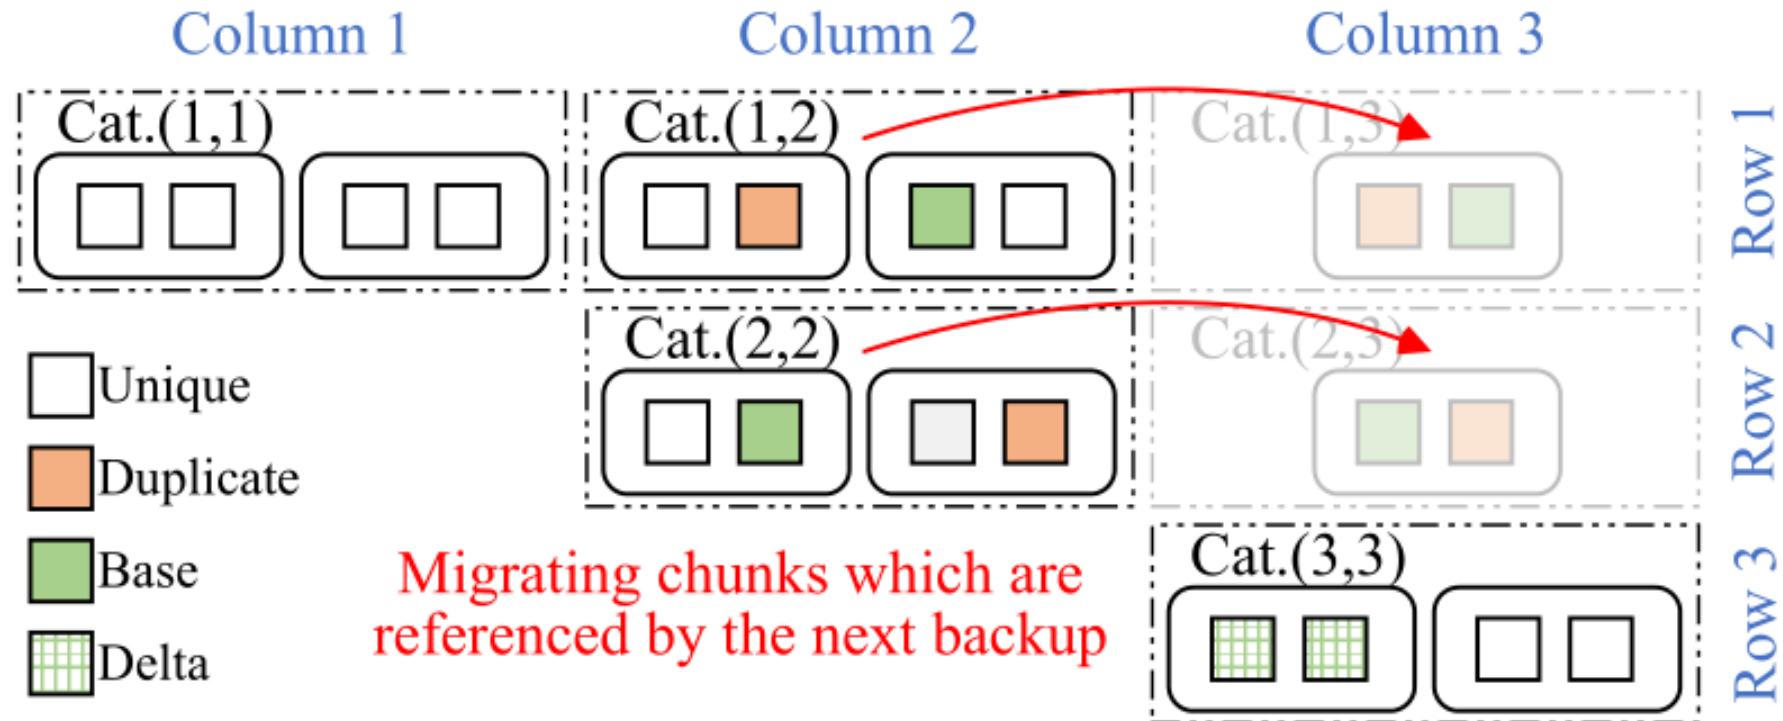


- NC_Backup1: A, B, C, D, E, F, G
- NC_Backup2: A, B, $\Delta(B')$, C, $\Delta(C')$, H, I, F, $\Delta(F')$, G
- NC_Backup3: A, J, C, H, $\Delta(H')$, I, F, $\Delta(F')$, G, $\Delta(G')$



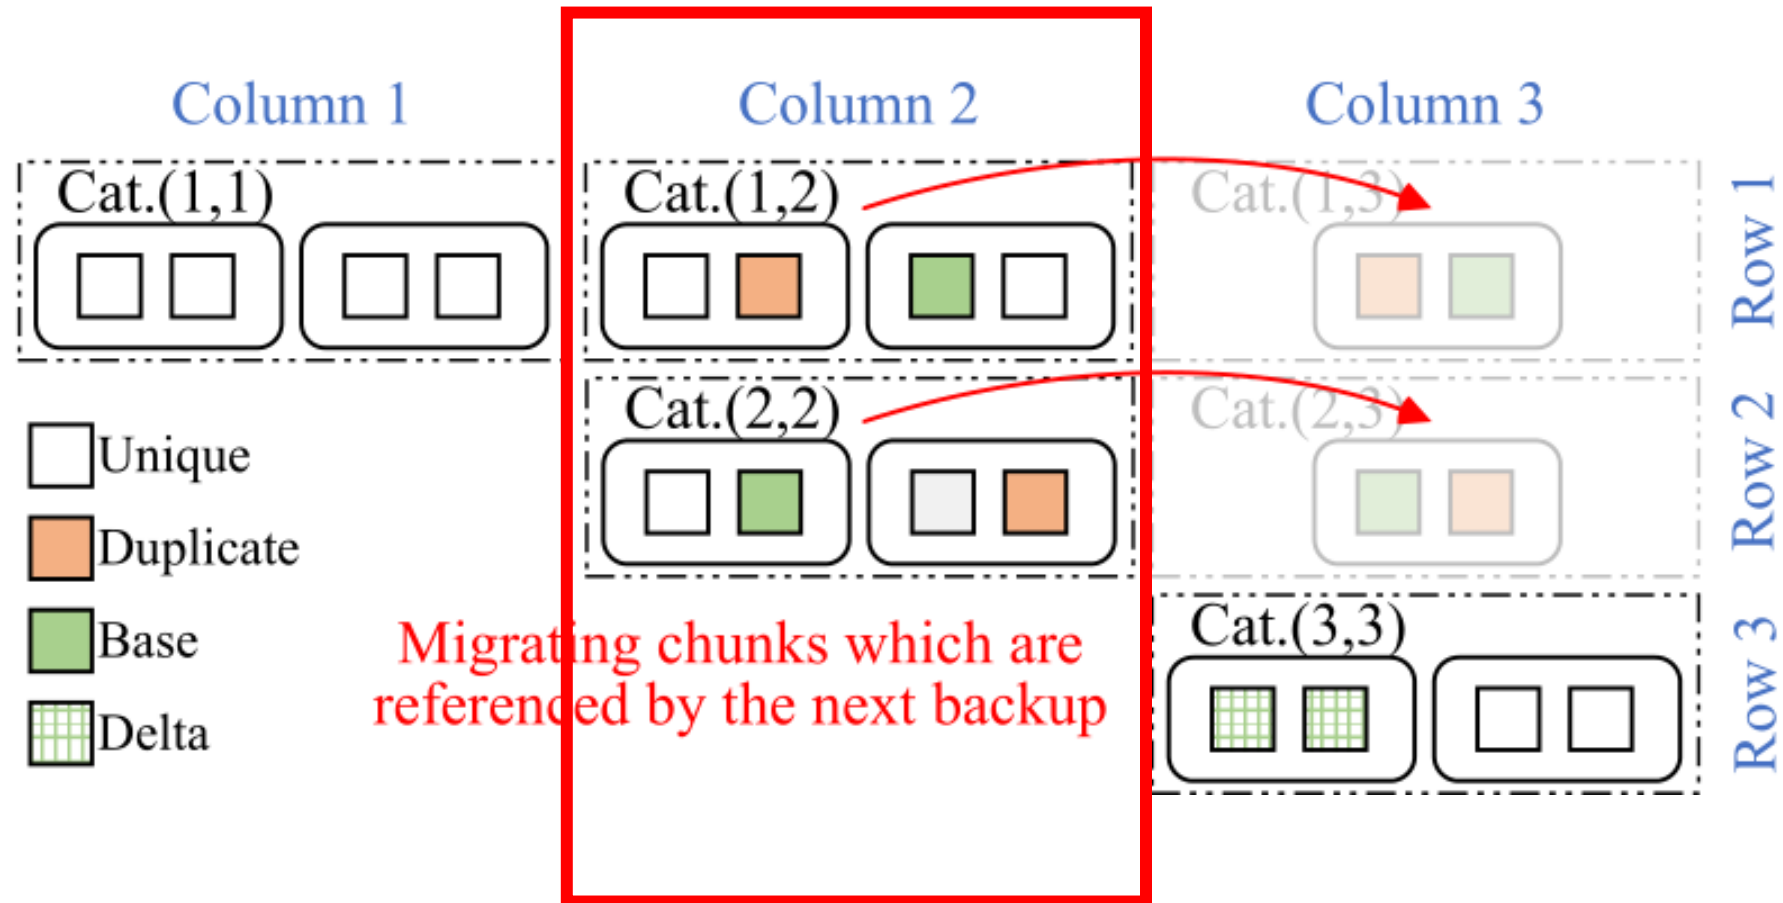
Design Delta-friendly Data Layout

- **Key Idea: Put the chunks from the same backup together**



Design Delta-friendly Data Layout

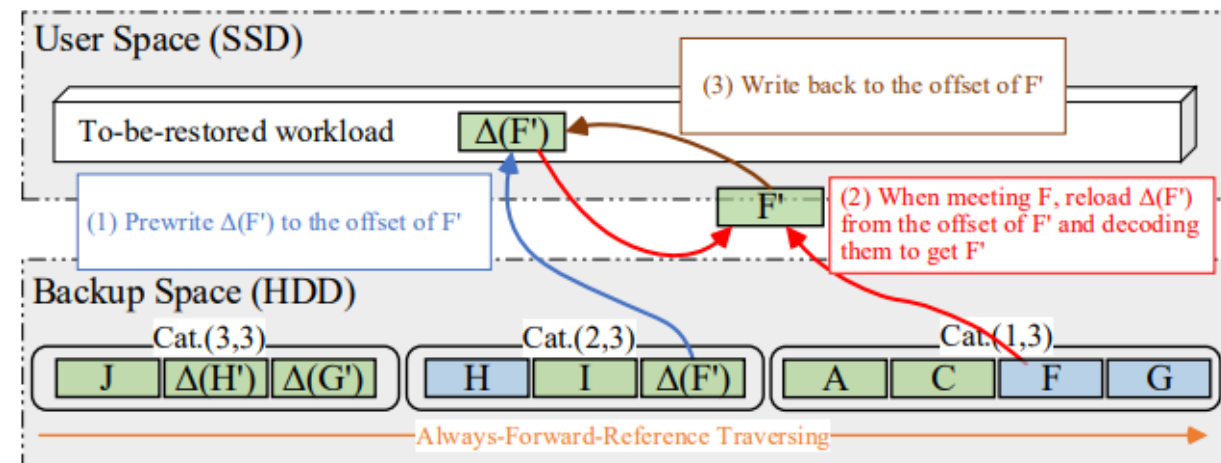
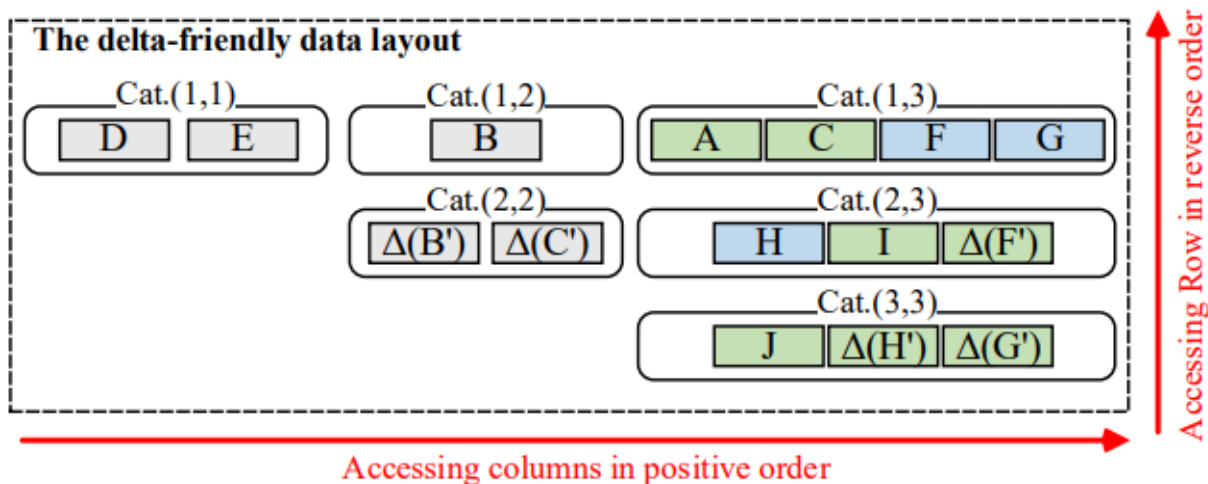
- **Key Idea: Put the chunks from the same backup together**

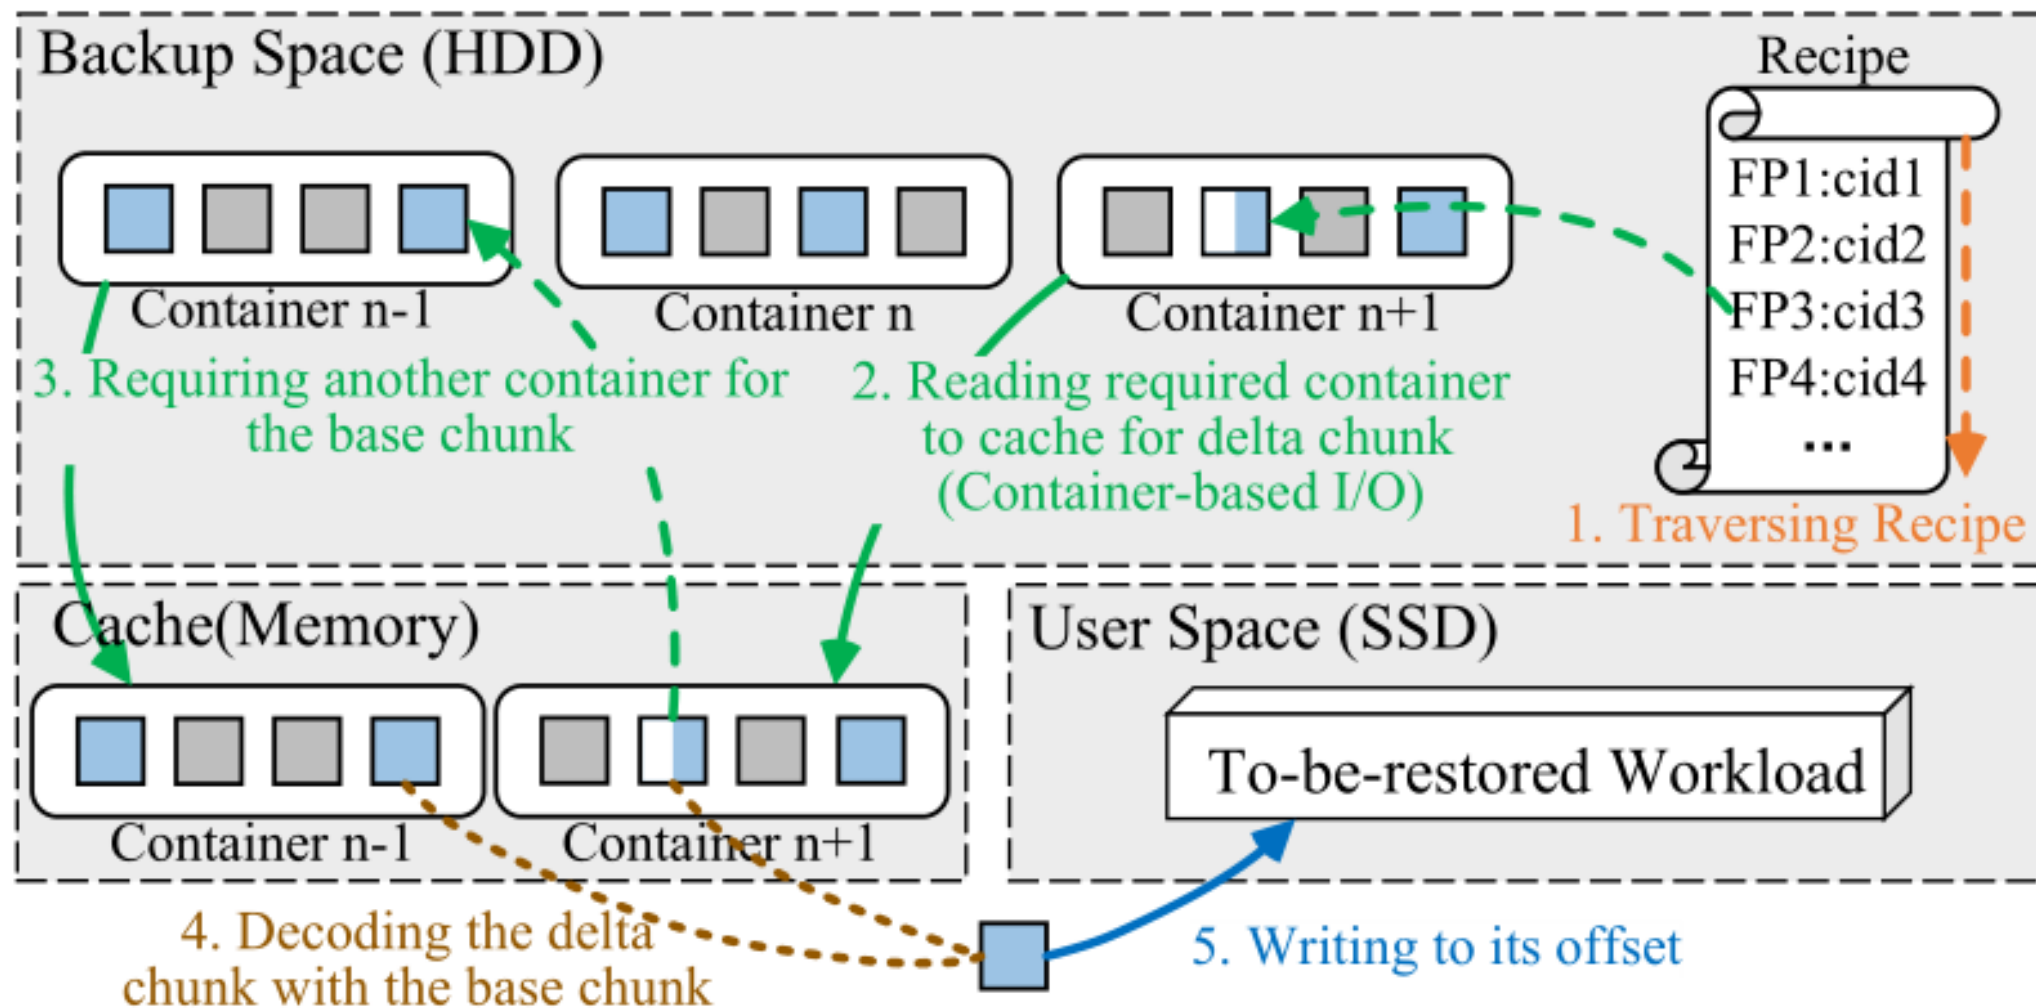


Design Always-Forward-Reference Traversing and Delta Prewriting

➤ Key Idea: Read the Container once and delta chunks before the base chunks

- A special path to traverse the restore-required chunks
 - Promises that delta chunks always appear before their base chunks
 - Rules to achieve AFR traversing
- Prewriting delta chunks
 - Asymmetric I/O characteristics of backup's/user's storage media
- Avoids repeatedly accessing restore-required chunks/containers





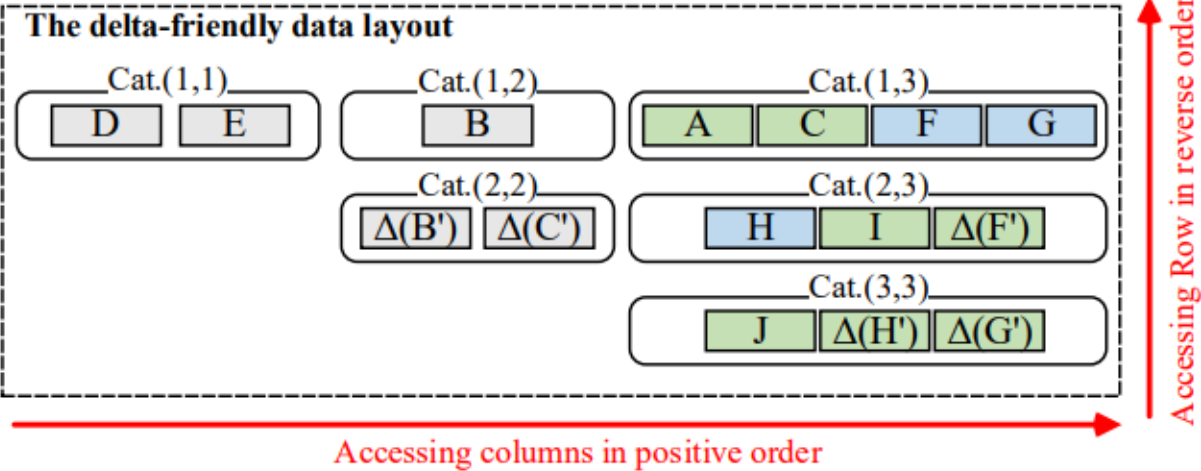
Design

Always-Forward-Reference Traversing and Delta Prewriting

➤ **Key Idea: Read the Container once and delta chunks before the base chunks**

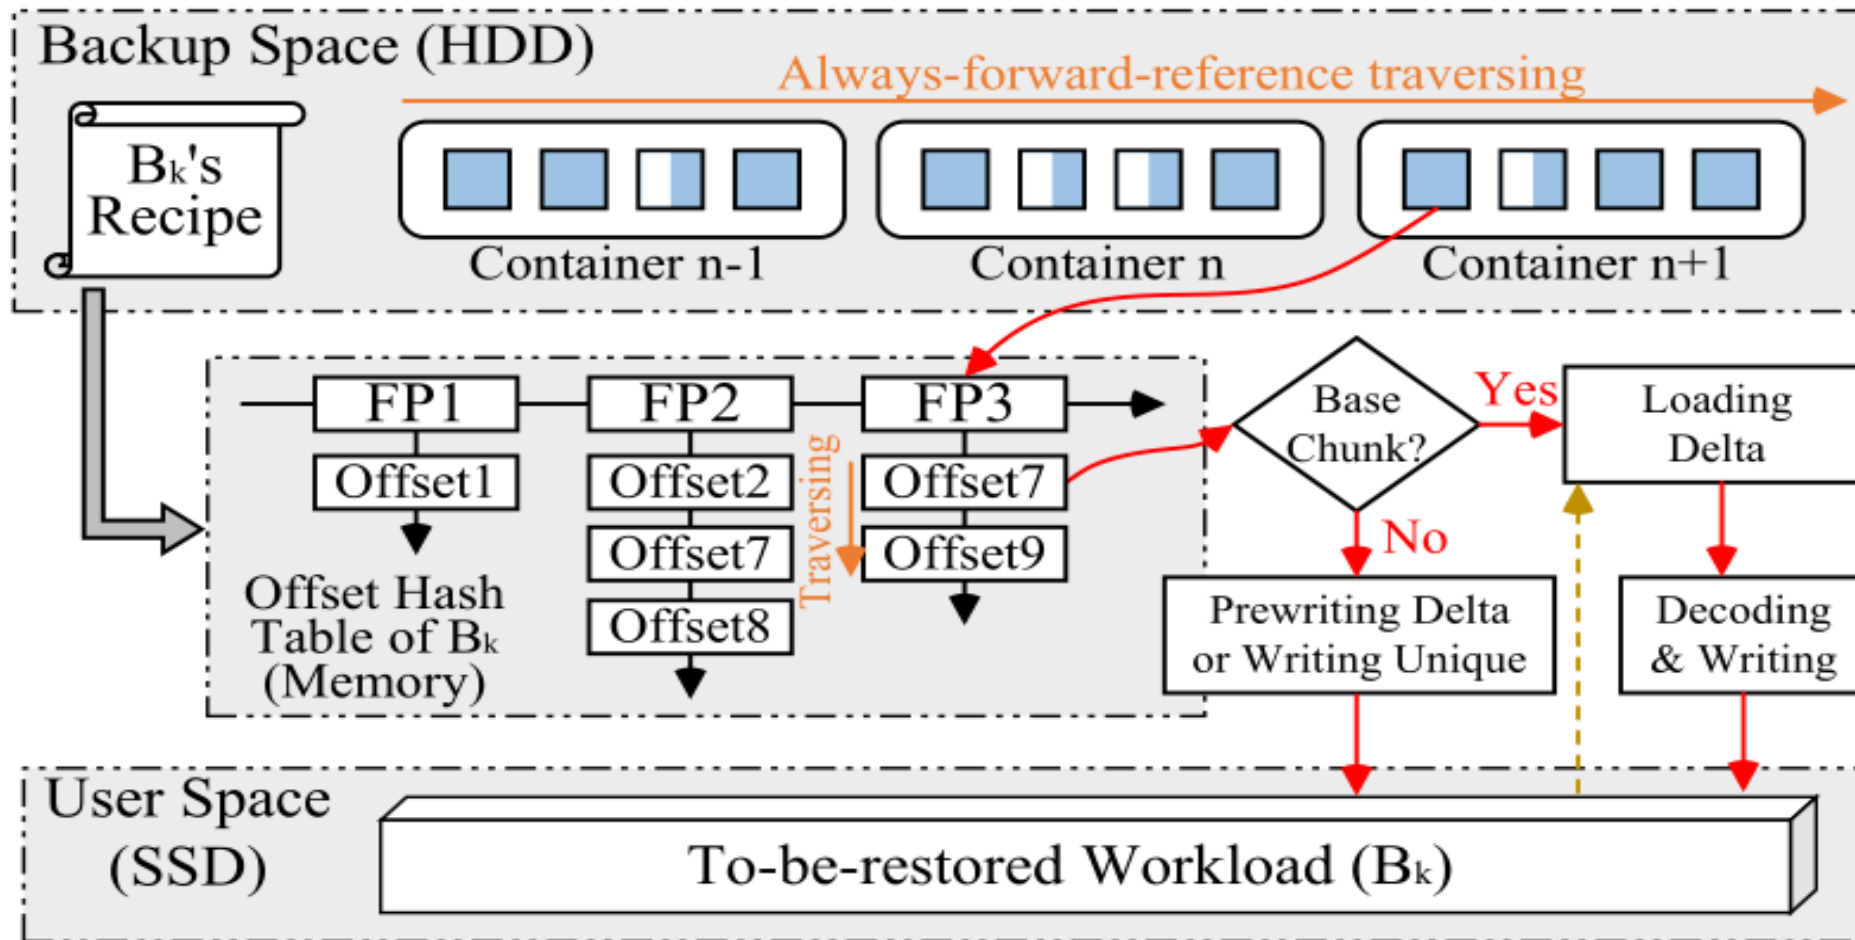
- A special path to traverse the restore-required chunks
 - Promises that delta chunks always appear before their base chunks
 - Rules to achieve AFR traversing
- Prewriting delta chunks
 - Asymmetric I/O characteristics of backup's/user's storage media
- Avoids repeatedly accessing restore-required chunks/containers

Delta Chunks' Positions		Corresponding Base Chunks' Possible Positions
Cat.(1,2)	⇒	Cat.(1,2), Cat.(1,3)
Cat.(2,2)	⇒	Cat.(1,2), Cat.(2,2), Cat.(1,3), Cat.(2,3)
Cat.(1,3)	⇒	Cat.(1,3)
Cat.(2,3)	⇒	Cat.(1,3), Cat.(2,3)



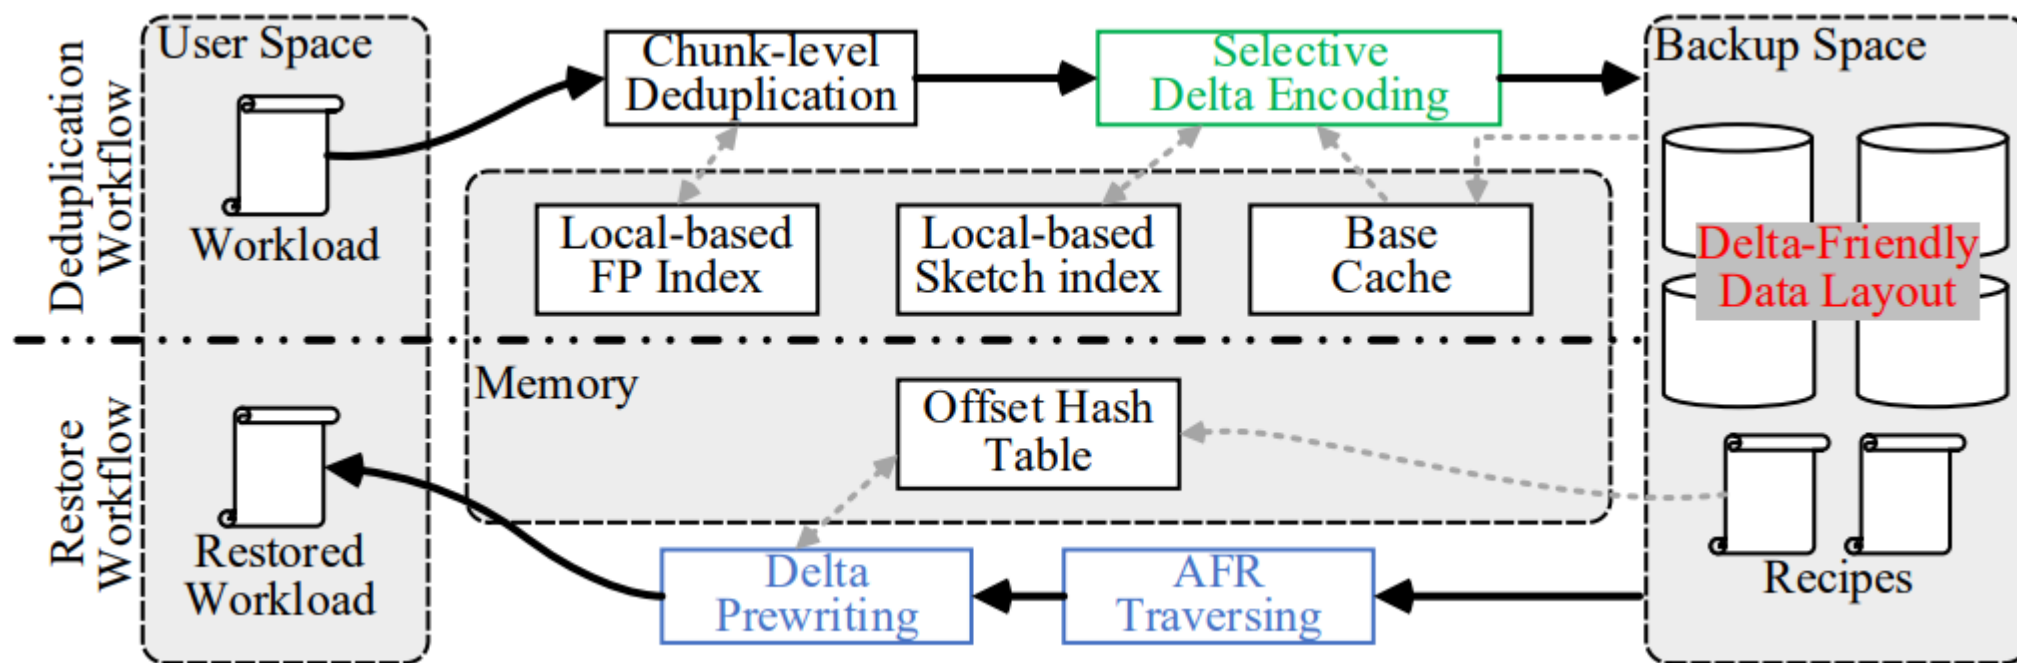
Design Always-Forward-Reference Traversing and Delta Prewriting

- **Key Idea: Read the delta chunks before the base chunks**



Architecture

- Techniques to address these three additional locality issues
 - Selective Delta Encoding
 - Delta-friendly Data Layout
 - Always-Forward-Reference Traversing and Delta Prewriting
- A fine-grained deduplication framework – MeGA



Evaluation

Experimental Setup


➤ Evaluated approaches

- **MeGA** Our proposed approach, using the three proposed techniques
- **Greedy** A fine-grained dedup approach with a greedy strategy
- **FGD** A fine-grained dedup approach with the Capping rewriting technique
- **CLD** A chunk-level dedup approach with Capping rewrite technique
- **MFD** A chunk-level ~~Greedy~~ dedup approach with an optimized data layout


➤ Dataset

Name	Original Size	Versions	Workload Descriptions
WEB	269 GB	100	Backups of website: news.sina.com, captured from Jun. to Sep. in 2016
CHM	279 GB	100	Source codes of Chromium project from v82.0.4066 to v85.0.4165
SYN	1.38 TB	200	Synthetic backups by simulating file create/delete/modify operations
VMS	1.55 TB	100	Backups of an Ubuntu 12.04 Virtual Machine


Evaluation backup speed



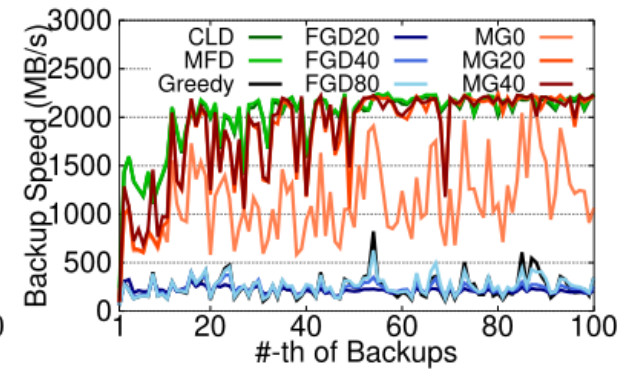
(a) WEB Dataset



(b) CHM Dataset



(c) SYN Dataset

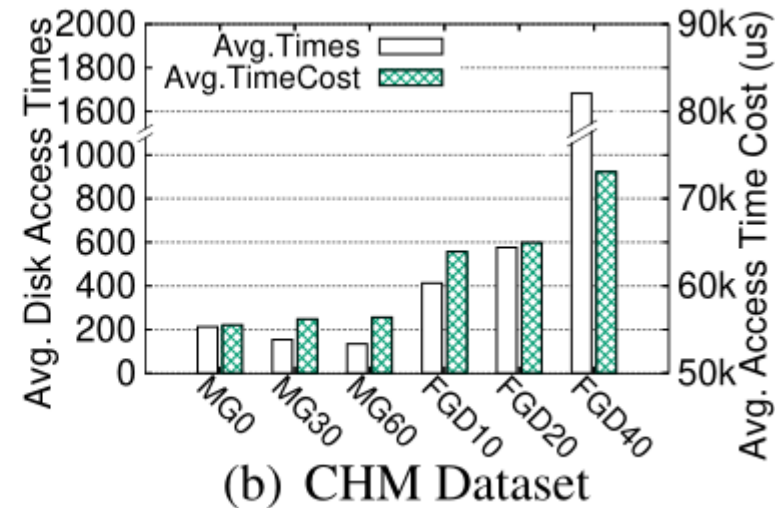



(d) VMS Dataset

- Applying several parameters for FGD and MeGA
- MeGA achieves a $4.47\text{--}34.45\times$ higher backup speed than Greedy

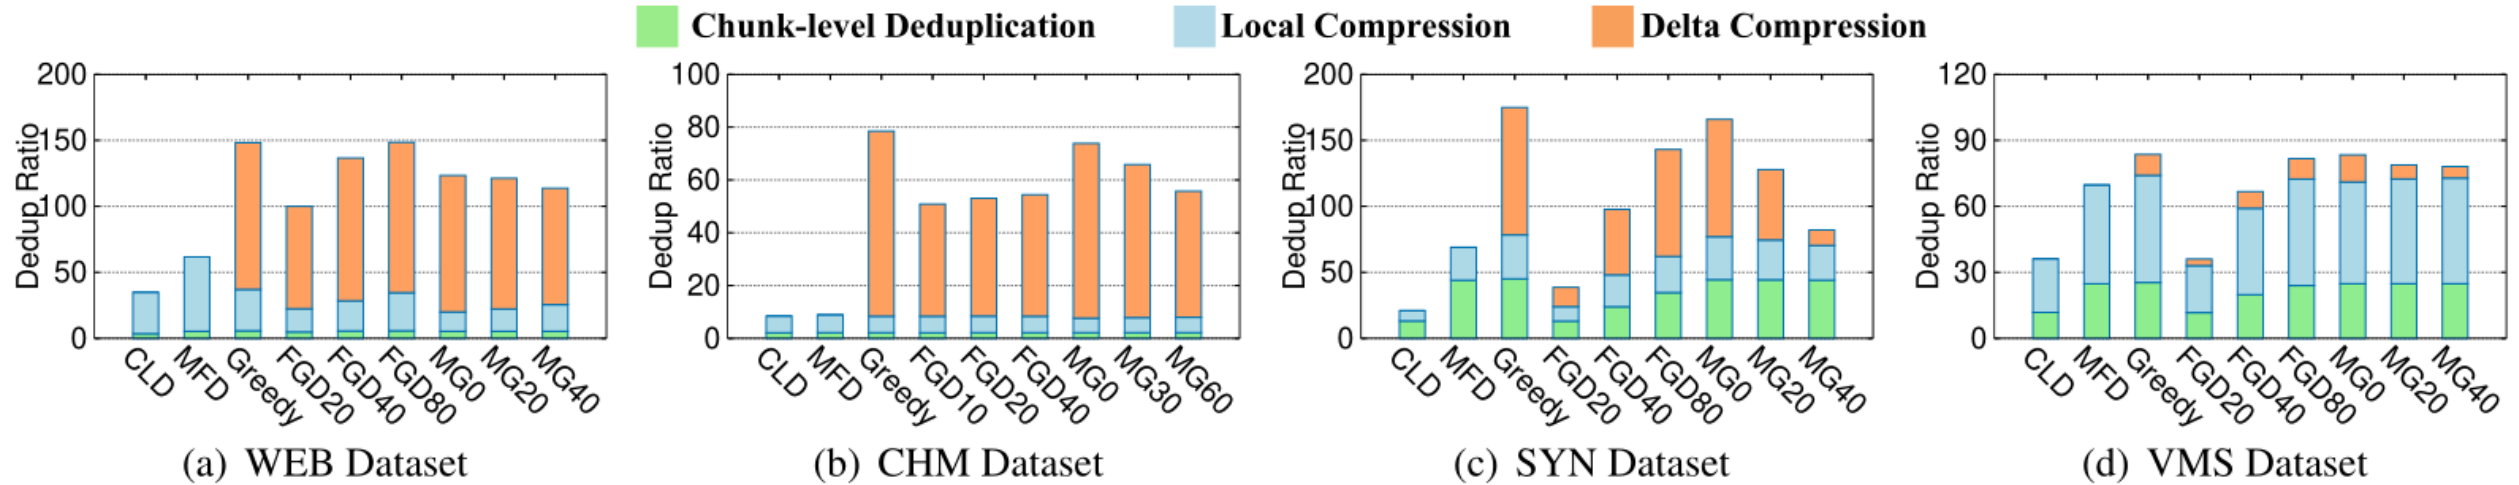
Evaluation

statistics about accessing disks for reading bases




- Selective Delta Encoding hugely reduces disk accessing times
- Skipping more delta encoding will lead to a better speed.

Evaluation Deduplication Ratio




- Fine-grained dedup achieves higher dedup ratio on most datasets
- There are few similar chunks in the VMS dataset
- MeGA preserves deduplication ratio advantage
- The deduplication ratio loss caused by Selective Delta Encoding is limited


Evaluation restore speed




(a) WEB Dataset



(b) CHM Dataset




(c) SYN Dataset




(d) VMS Dataset

- MeGA achieves a 30–105× higher restore speed than Greedy

Evaluation statistics about accessing disks for required chunks



(a) WEB Dataset



(b) CHM Dataset

- Our data layout hugely reduces the restore-involved containers
- Always-Forward-Reference Traversing and Delta Prewriting avoid the repeatedly accessing

Conclusion

