

# **ZoneAlloy: Elastic Data and Space Management for Hybrid SMR Drives**


Fenggang Wu Bingzhe Li Zhichao Cao Baoquan Zhang  
Ming-Hong Yang Hao Wen David H.C. Du

**HotStorage 19**

# SMR HDDs

## ➤ SMR structure

- Tracks are overlapped in a shingled fashion to achieve higher areal density on the same disk platter
- Track width is **smaller** than the magnetic reader head → need **Read-Modify-Write** to update blocks in-place




## ➤ The Problem: SMR Update Overhead

# SMR, CMR and H-SMR HDDs

Conventional Writing



Shingled Writing




## ➤ CMR:

- Low capacity; High performance

## ➤ SMR:

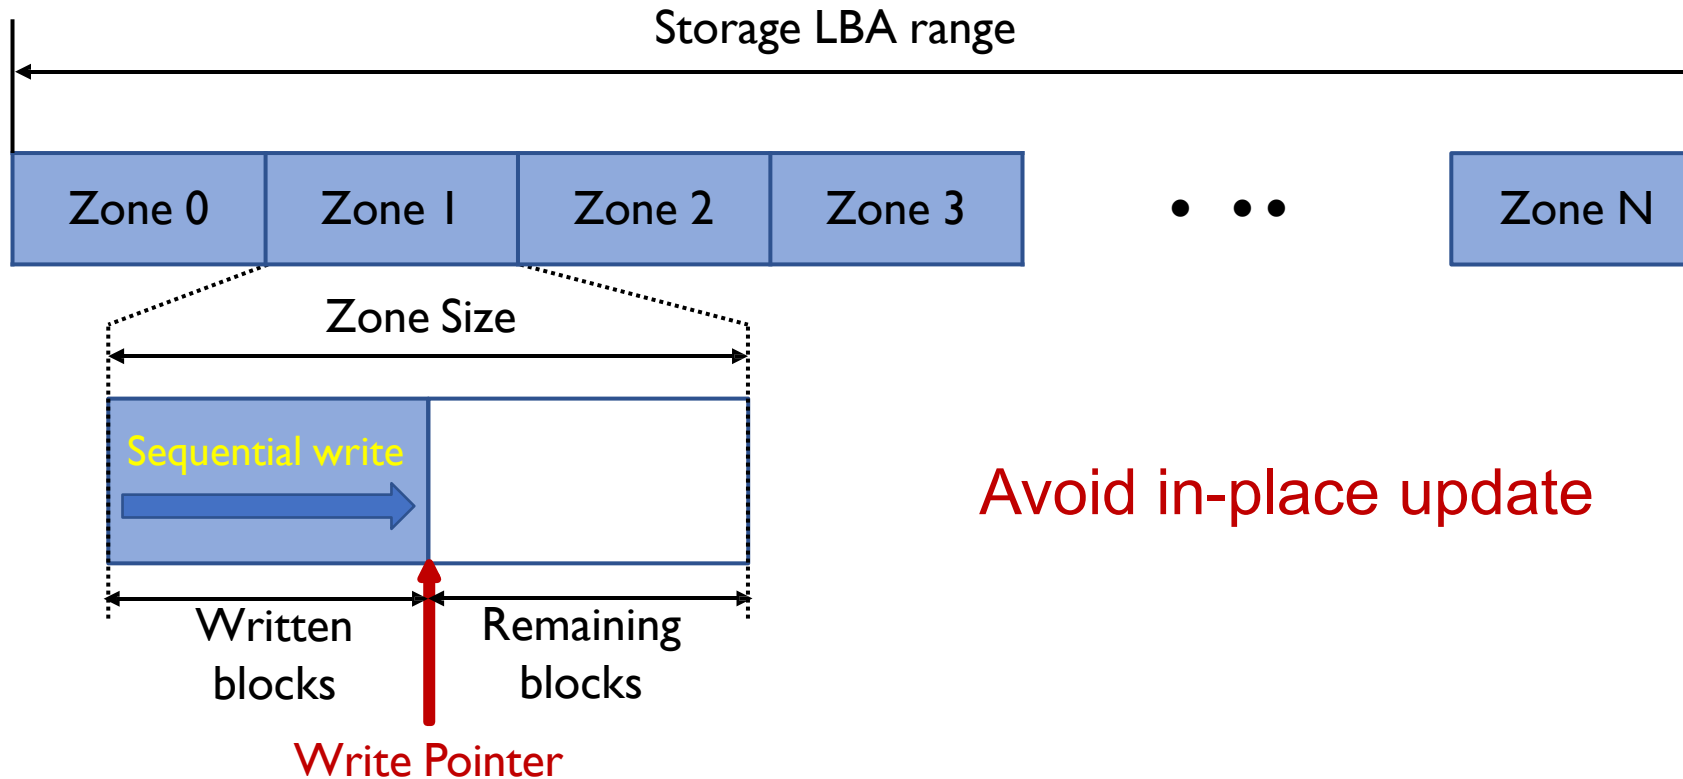
- High capacity; Low performance
- High update overhead (latency)

- H-SMR manages the **performance/capacity trade-off** on the disk



# Zone Structure for H-SMR

- “Zone” is a consecutive LBA space with a size of 256MiB
  - **SMR zone** has a write pointer indicating where the next write should go to enforce sequential write.

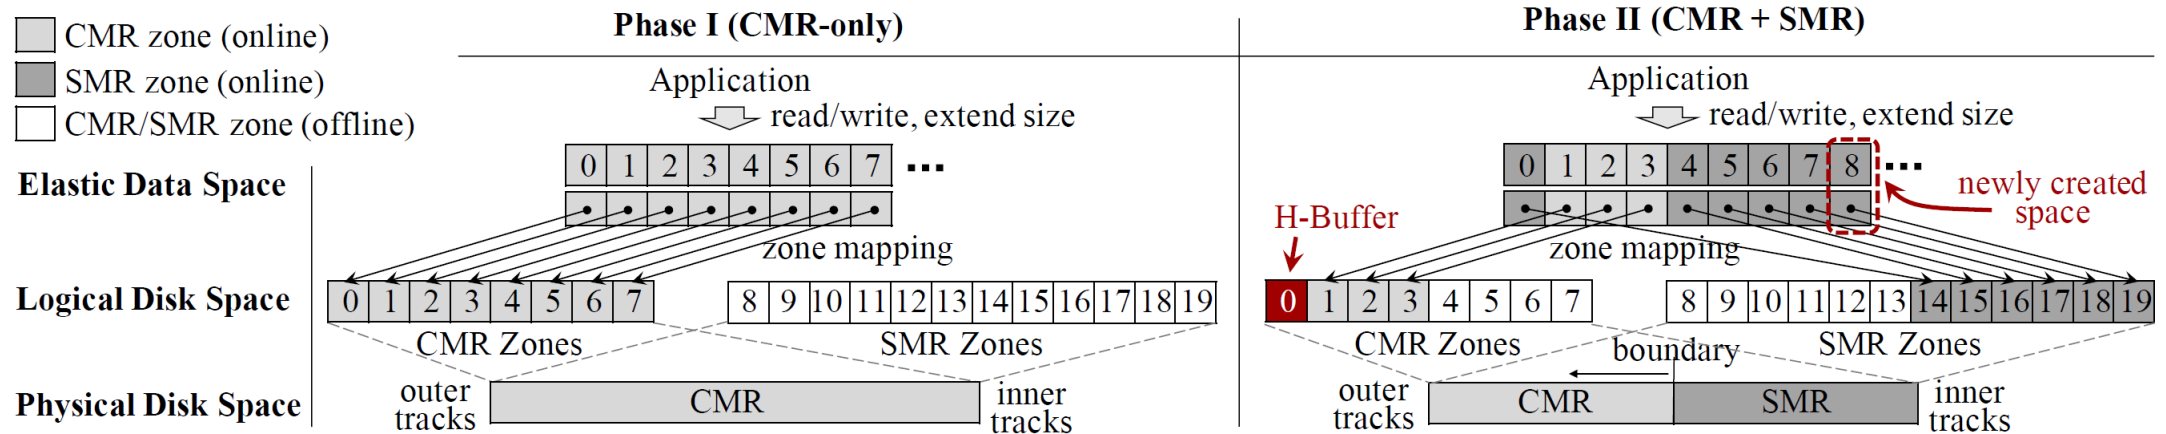


# Four Problems with H-SMR

- How to arrange the format layout and place data accordingly?
  - Coexistence of CMR and SMR
- How to perform format conversion efficiently?
  - SMR density is about 1.5 times that of CMR
- How to reduce SMR update overhead?
  - In-place update SMR data blocks
- How to adapt to dynamic workloads?
  - “Hot” and “Cold” data requires different latency

# Zone Mapping and Two-phase Allocation

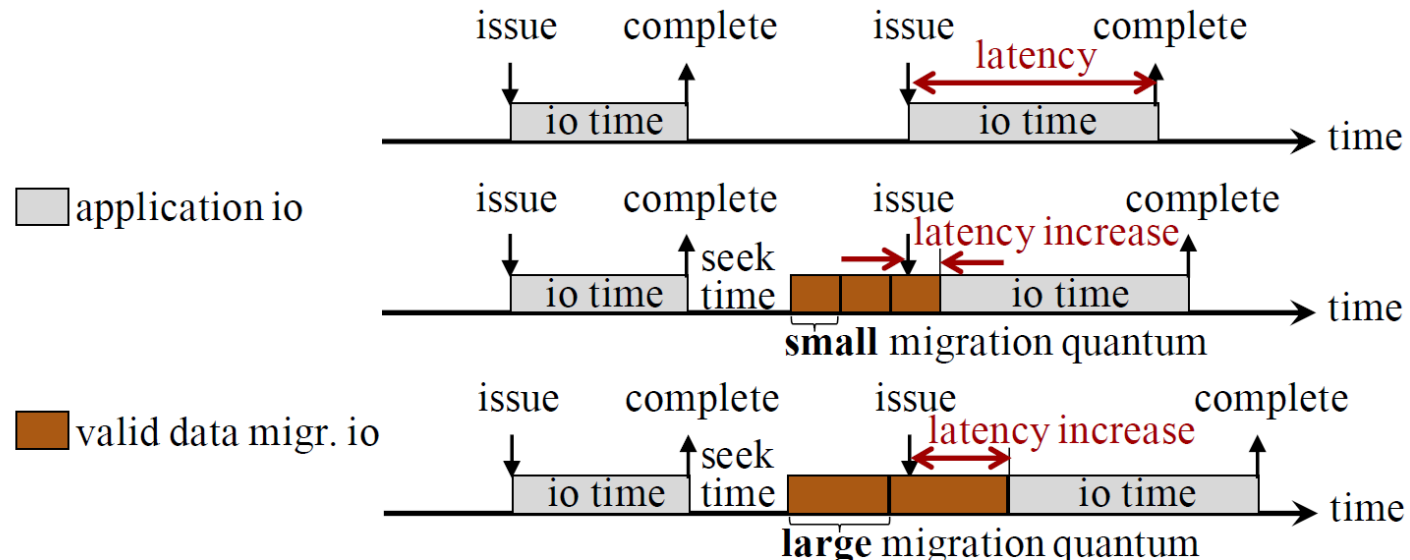
- H-SMR divide space into zones (CMR+SMR)
- H-SMR provide elastic data space
  - Full CMR (phase I) to CMR+SMR (phase II)



HDDs outer tracks have higher performance than inner tracks

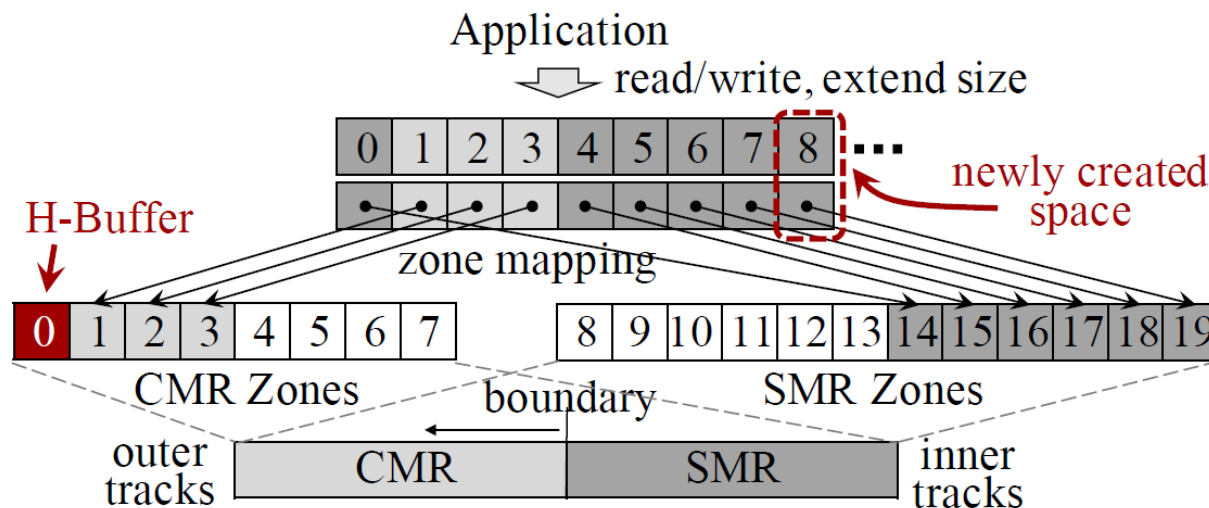
# Quantized Migration

- Larger migration quantum →
  - conversion finish sooner
  - higher latency increase
- Application specifies the acceptable increase of latency



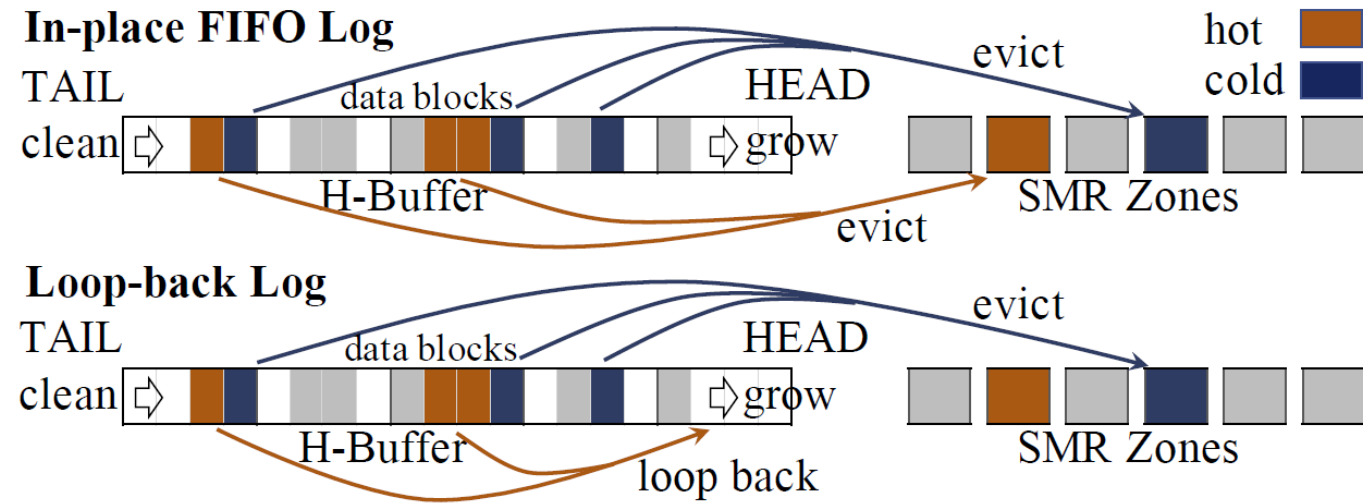
# H-Buffer based SMR Update

- H-Buffer: small host-controlled CMR cache
  - Absorb and migrate SMR updates in batches
- Cache policy: LRU?
  - Block-based LRU leads to poor performance: zone update for each block
  - Zone-based LRU fragment free space of H-Buffer, causing random I/Os (Low performance)





# H-Buffer based SMR Update



## ➤ In-place FIFO Log


- Evict blocks of same zone together → reduce zone update
- Frequently updated data blocks will come back to H-Buffer soon

## ➤ Loop-back Log:

- Use LRU to predict “hot” zones, then move blocks back to HEAD

# Zone-Swap Scheme

- SMR swap candidates
  - Zones with an occupancy in H-Buffer(during last epoch) above a threshold
- CMR swap candidates
  - CMR zones that were not updated during last epoch



# Evaluation

## ➤ Platforms & Traces


- H-SMR simulator based on DiskSim.
- Microsoft Research (MSR) Cambridge traces

## ➤ Methodology

- SMR to CMR areal density ratio is set to 1.5:1
- Only the traces with a write footprint greater than 10GB are used

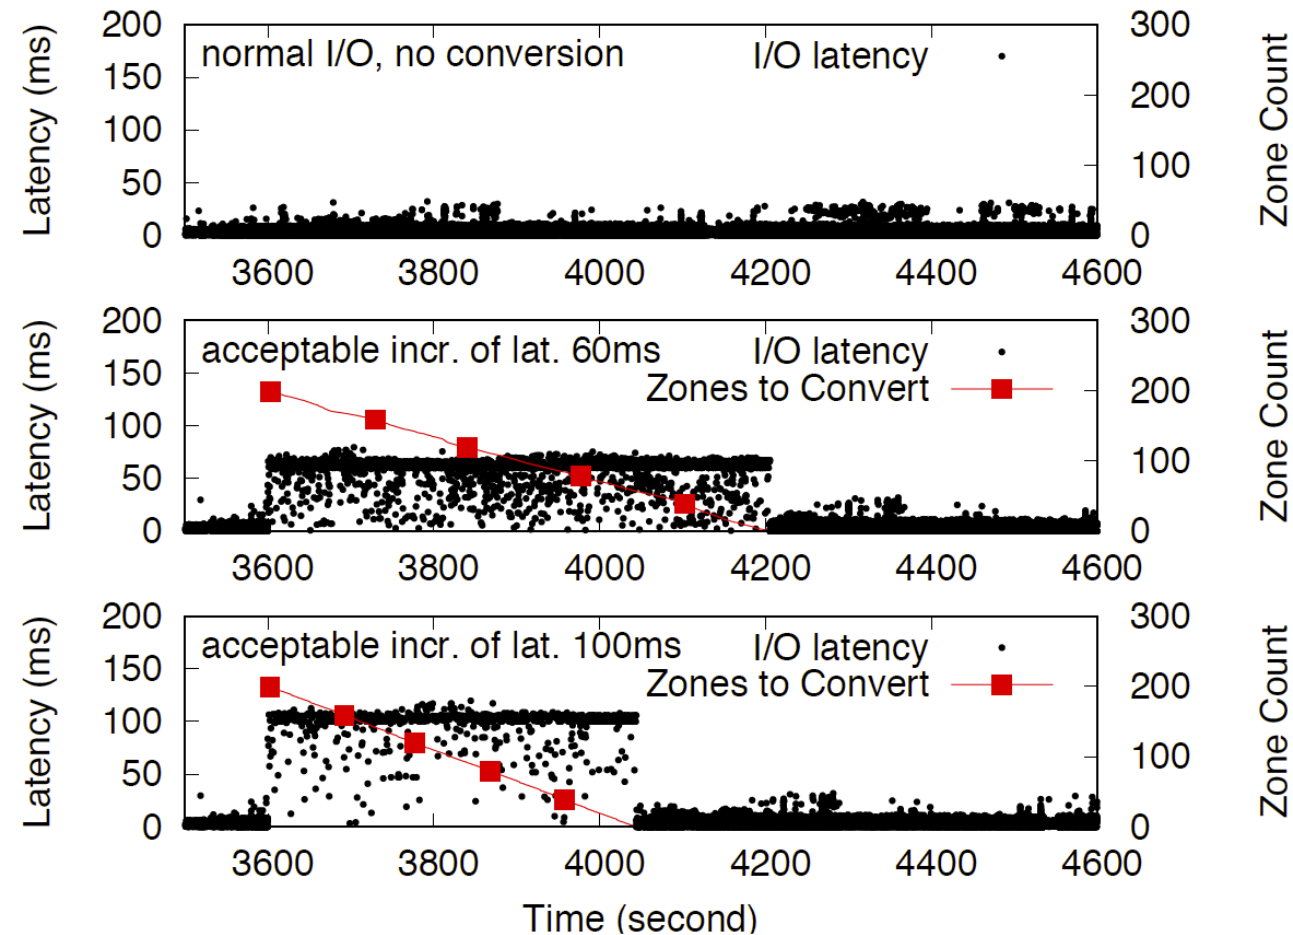
# Overall Latency

- H-Buffer to SMR partition size ratio is set to as low as 0.02%, and the disk usage is set to 99.9%
  - Improved-lru: zone-based LRU H-Buffer



Zone-alloy achieves lowest latency, and effect of zone-swap is significant in proj\_0/usr\_1 datasets

# Quantized Migration



➤ Application requests 25GiB of new space (200 zones to convert) with acceptable increase of latency set to 60 ms and 100 ms

- The higher the acceptance delay, the faster the conversion
- Delay not changed when no conversion

# Conclusion

- ZoneAlloy, an elastic data and space management scheme that hides the H-SMR details and presents an elastic data space to the application.
  - H-Buffer and Zone-swap to reduce SMR update overhead
  - Two-phase allocation to support elastic data space
  - Trade-off to perform conversion efficiently while bounding their performance degradation